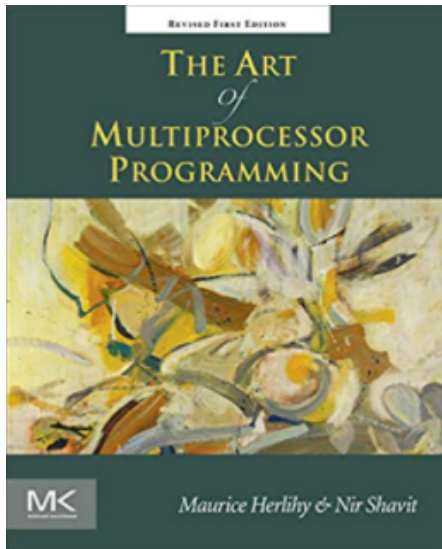


Programmation répartie

- » Carole Delporte
- » Moodle :
 - » IFECY 140 Programmation Répartie
- » programmation :java
- » Controle des connaissances: 50% TP et 50% exam

» Bibliographie



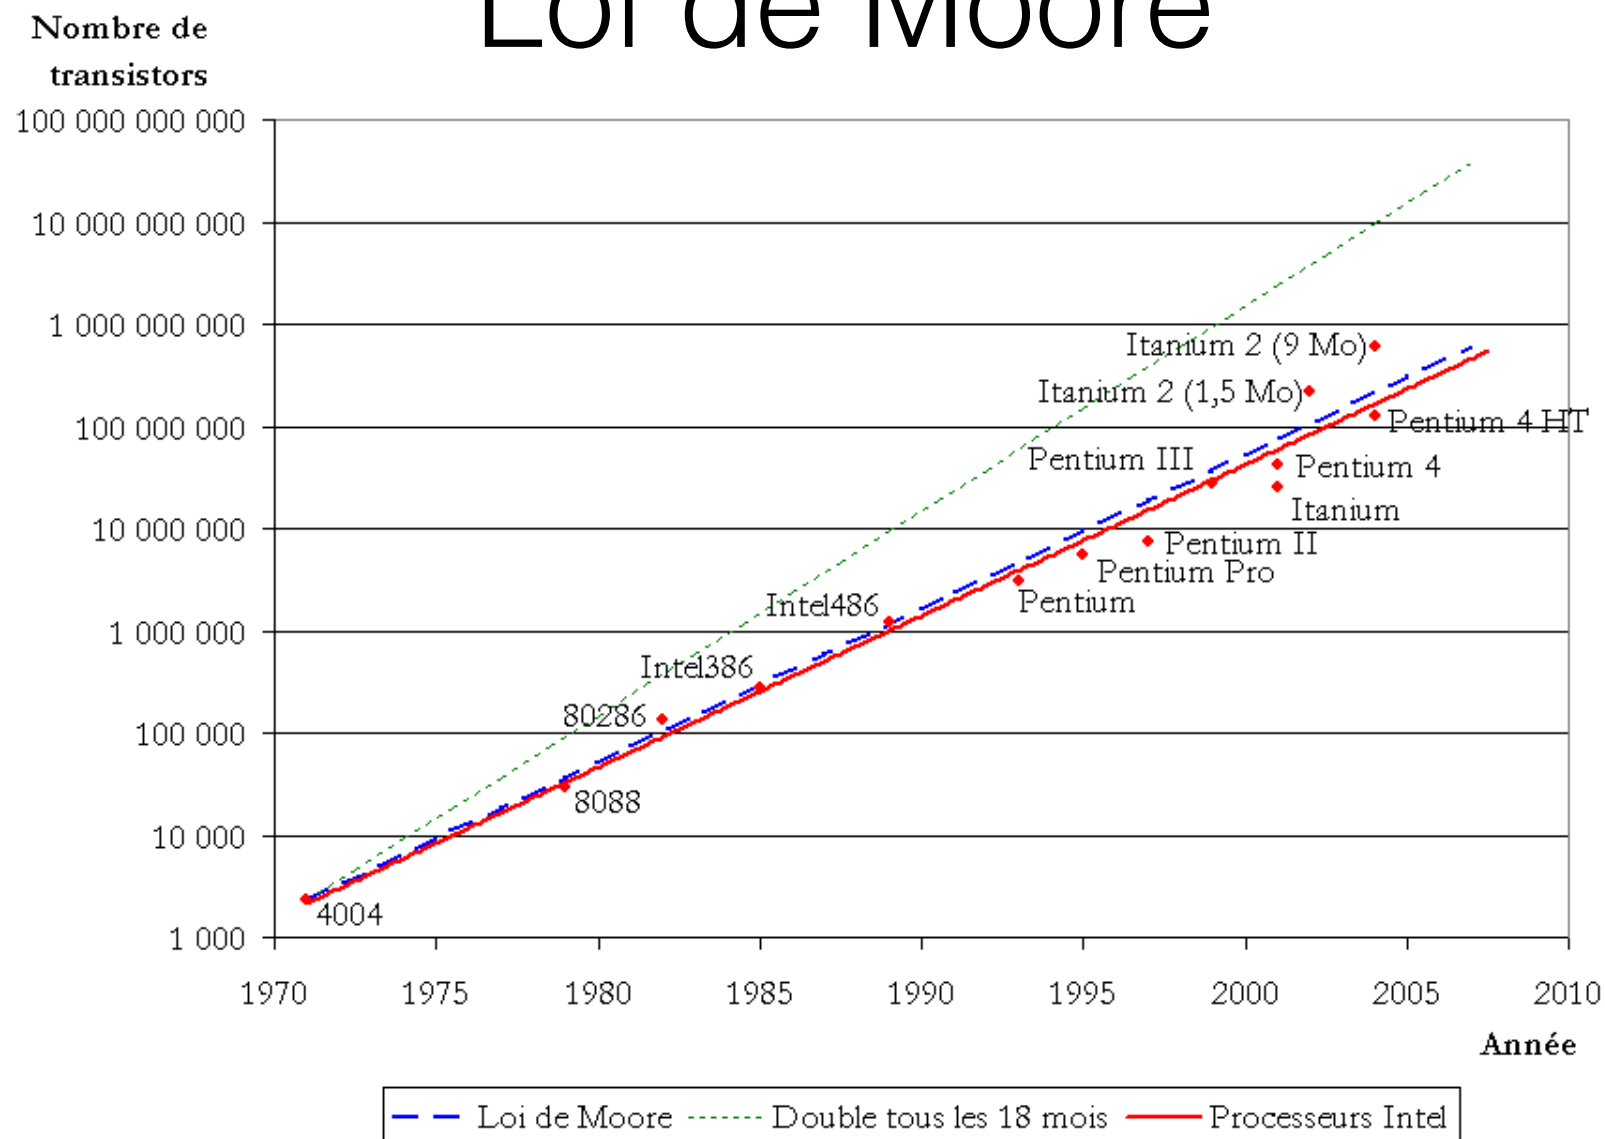
» The art of Multiprocessor programming

» Herlihy & Shavit

Companion slides for
The Art of Multiprocessor Programming
by Maurice Herlihy & Nir Shavit

Introduction

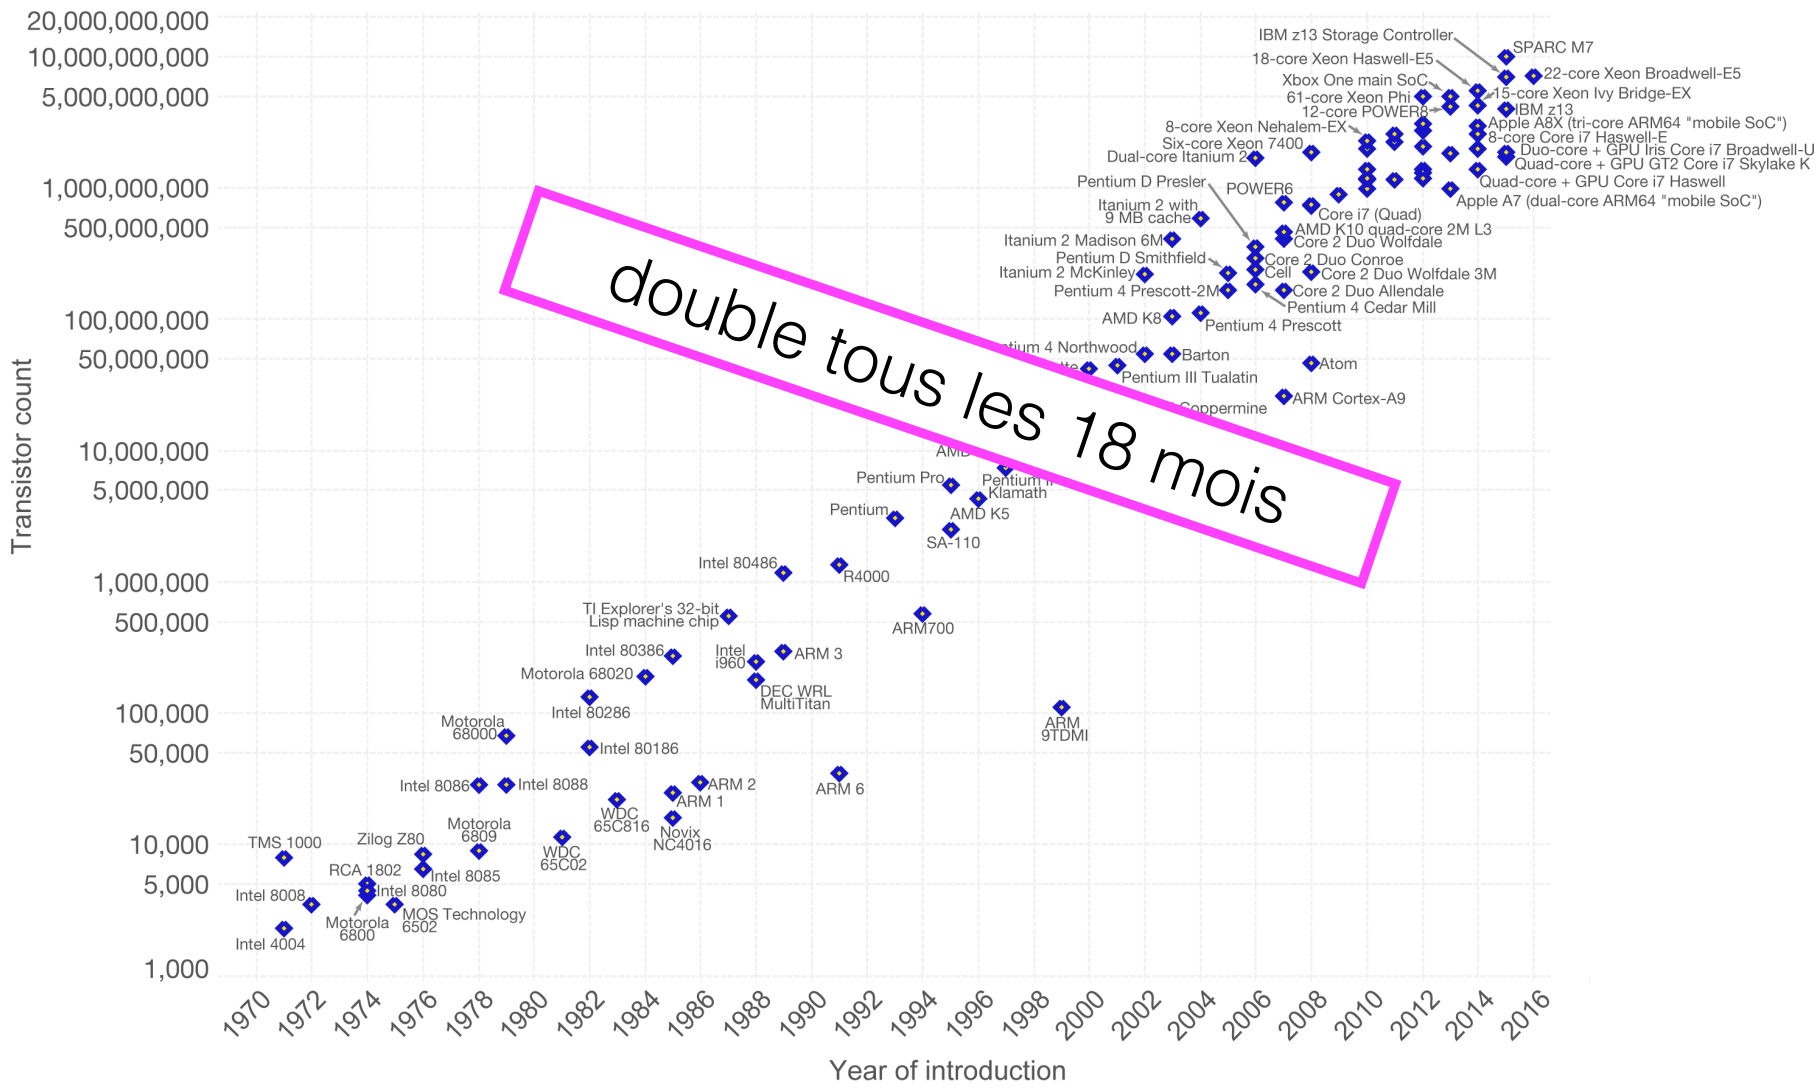
Loi de Moore



Moore's Law – The number of transistors on integrated circuit chips (1971-2016)

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years.

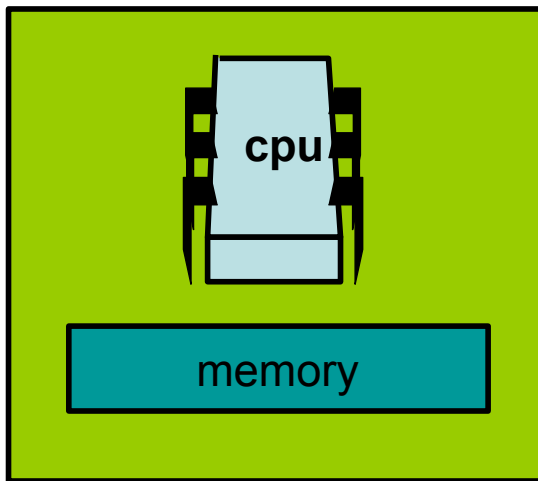
This advancement is important as other aspects of technological progress – such as processing speed or the price of electronic products – are strongly linked to Moore's law.



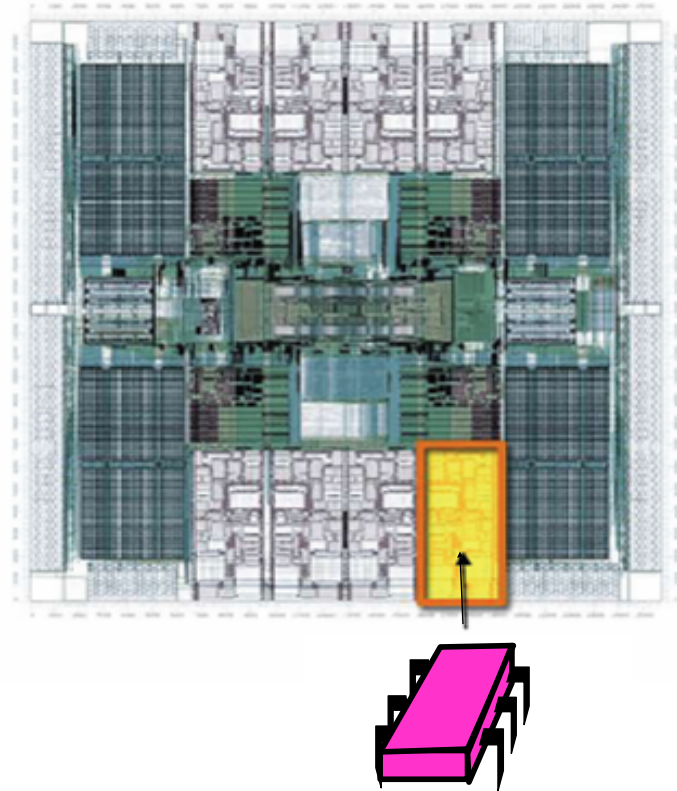
Data source: Wikipedia (https://en.wikipedia.org/wiki/Transistor_count)
The data visualization is available at [OurWorldinData.org](https://ourworldindata.org). There you find more visualizations and research on this topic.

Licensed under [CC-BY-SA](https://creativecommons.org/licenses/by-sa/4.0/) by the author Max Roser.

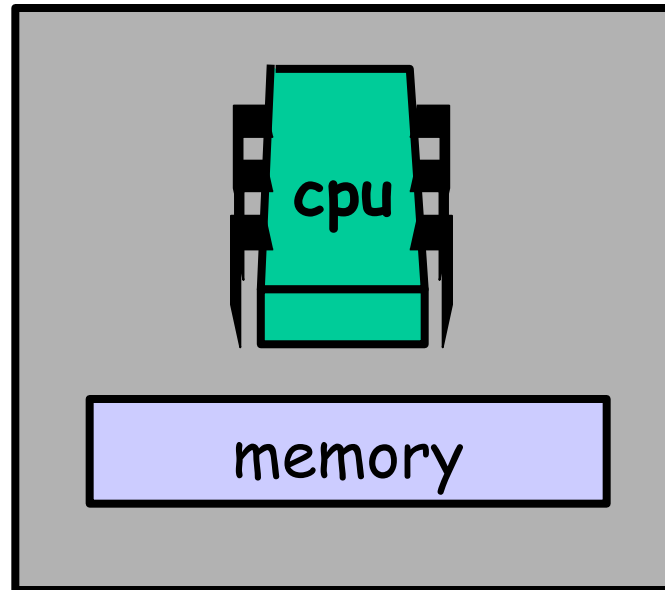
Uniprocasseur



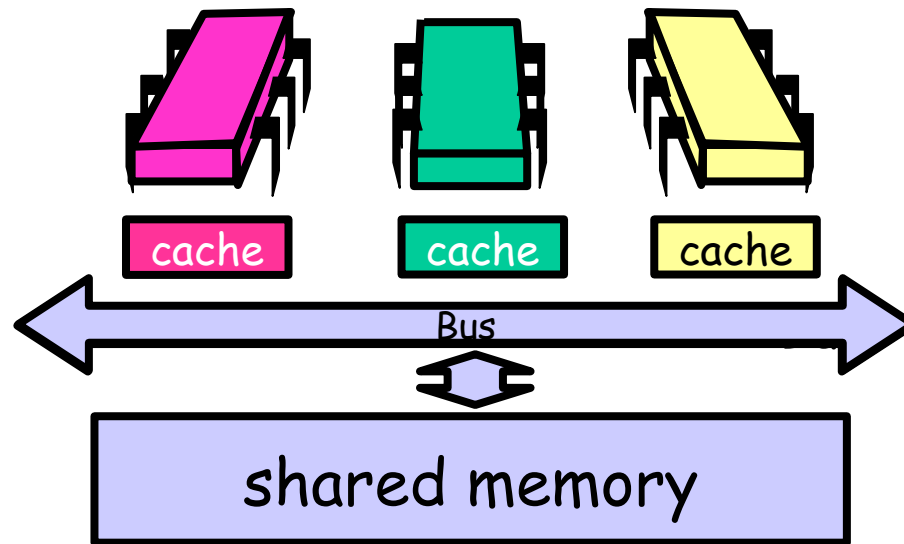
Mémoire partagée multicore



Still on some of your desktops: The Uniprocessor

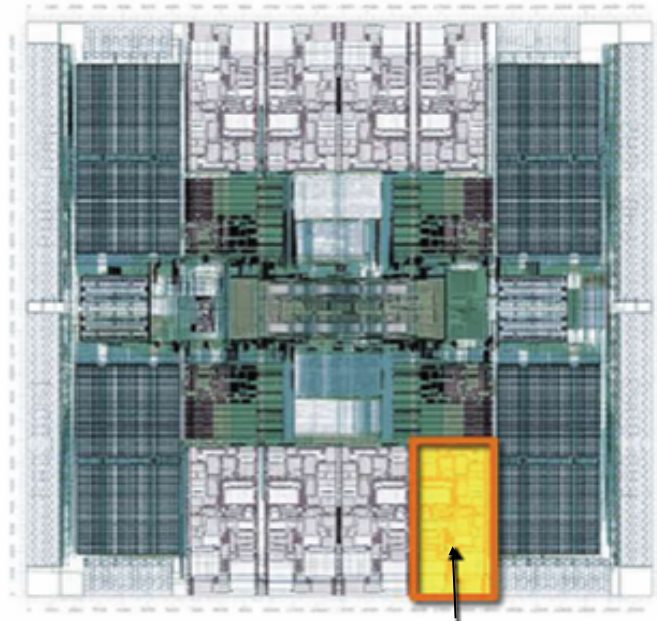


In the Enterprise:
The Shared Memory Multiprocessor
(SMP)

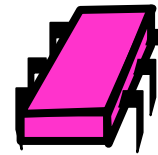


Your New Desktop:
The Multicore Processor
(CMP)

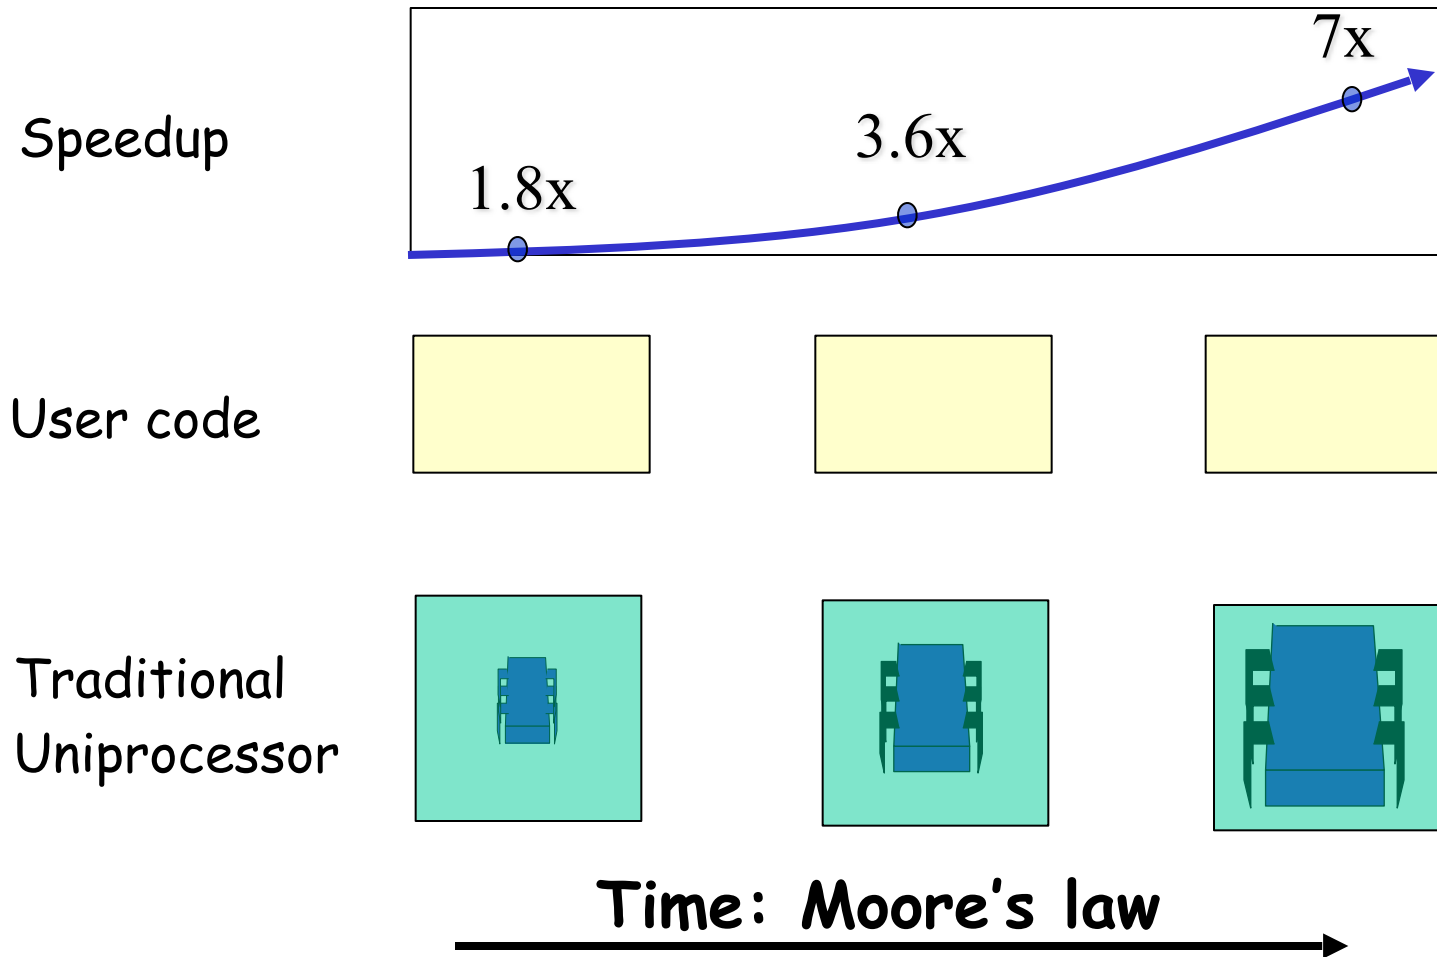
All on the
same chip



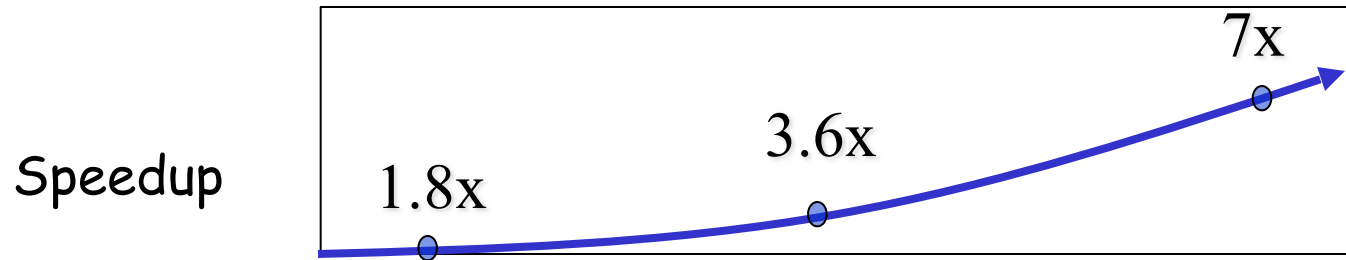
Sun
T2000
Niagara



Traditional Scaling Process

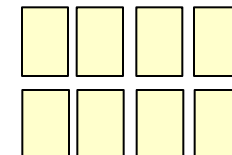
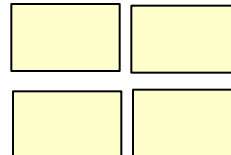


Multicore Scaling Process

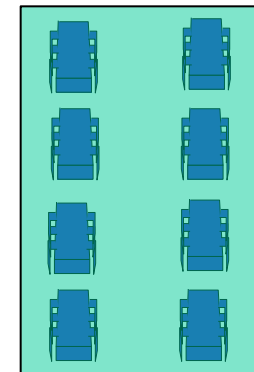
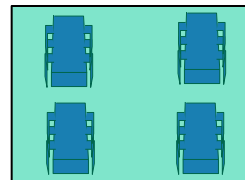


Theory

User code

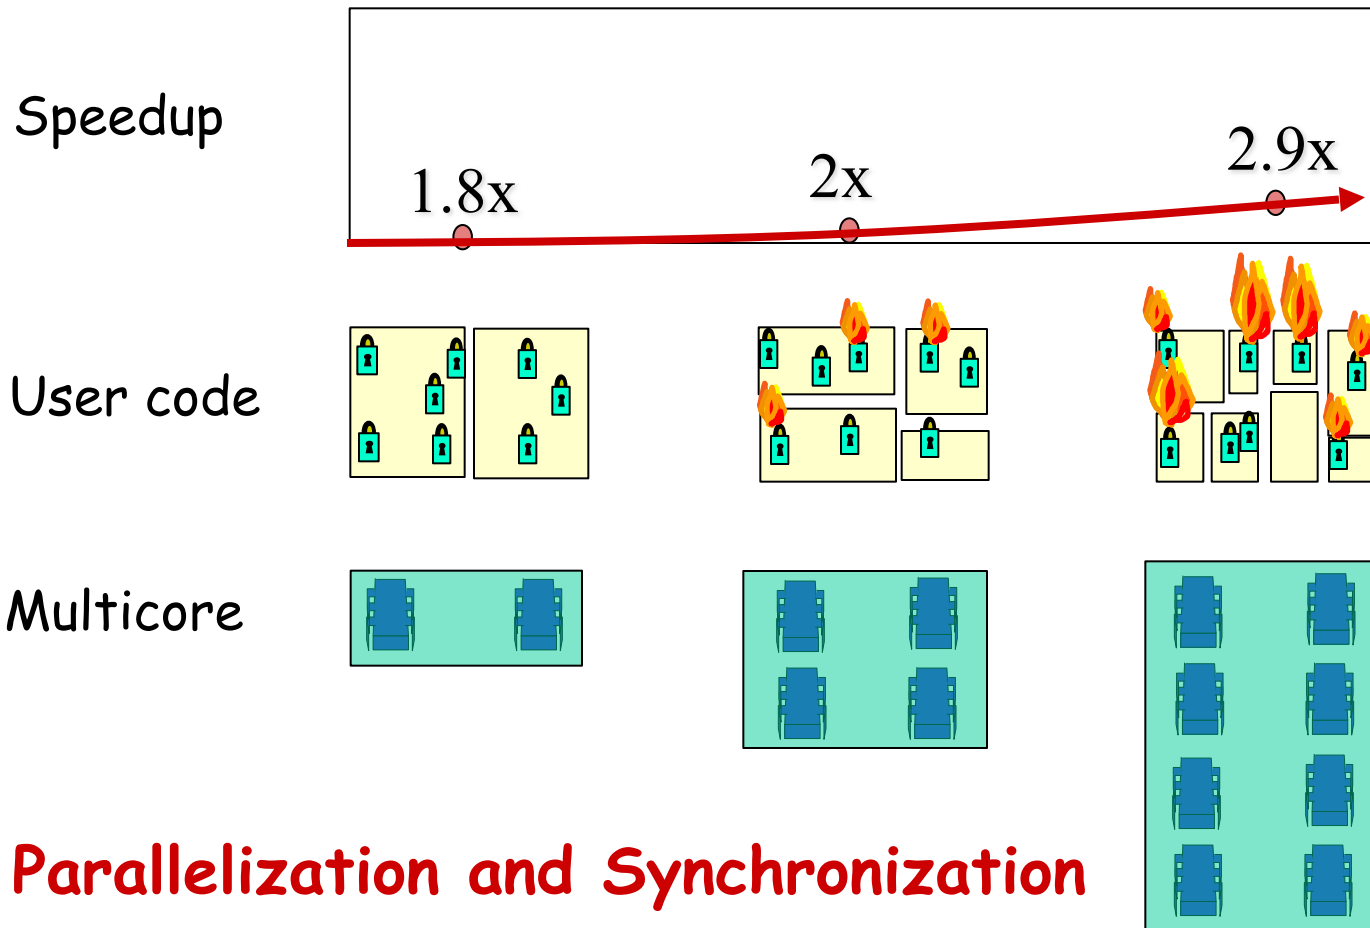


Multicore



Unfortunately, not so simple...

Real-World Scaling Process

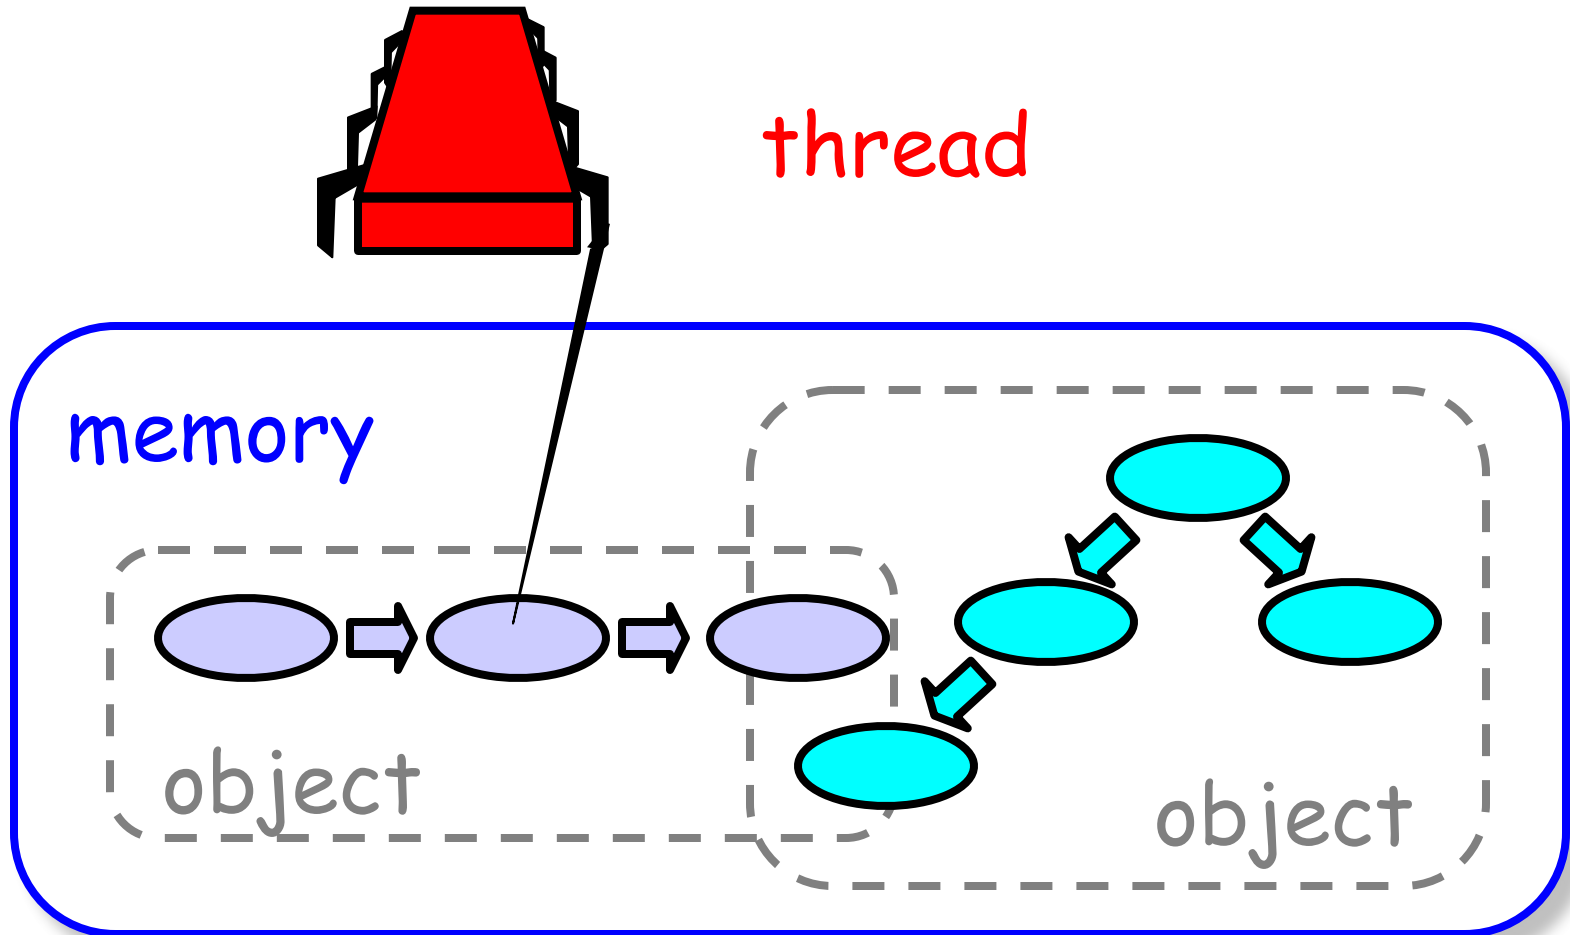


**Parallelization and Synchronization
require great care...**

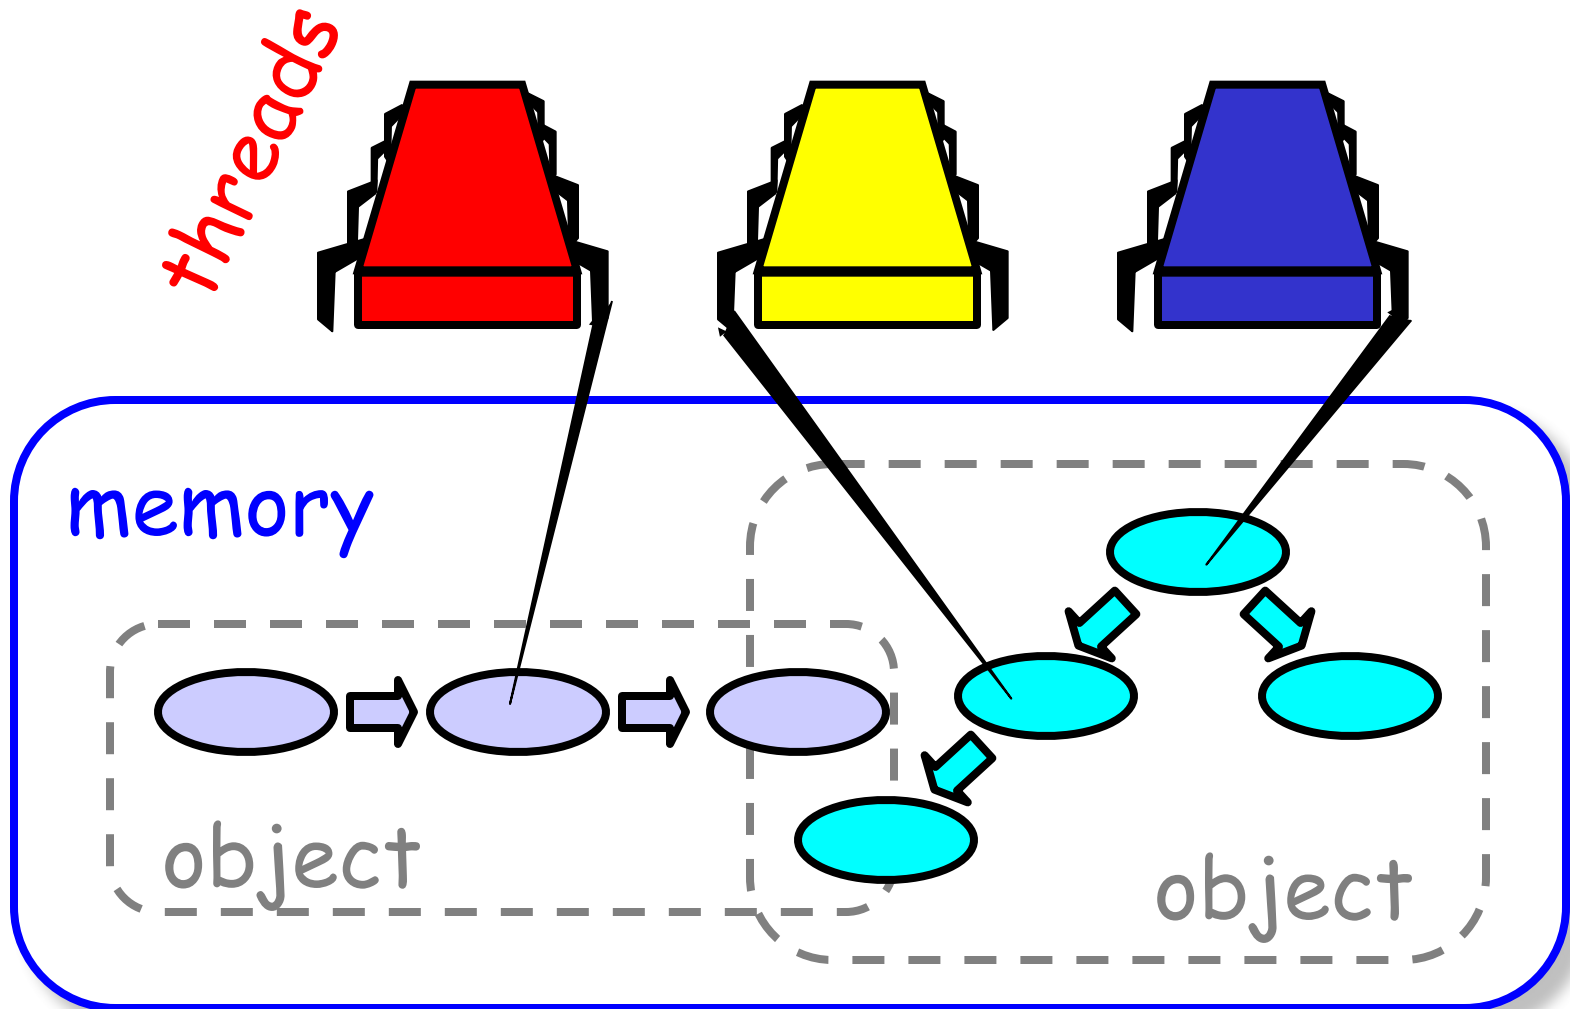
Multicore Programming: Course Overview

- Fundamentals
 - Models, algorithms, impossibility
- Real-World programming
 - Architectures
 - Techniques

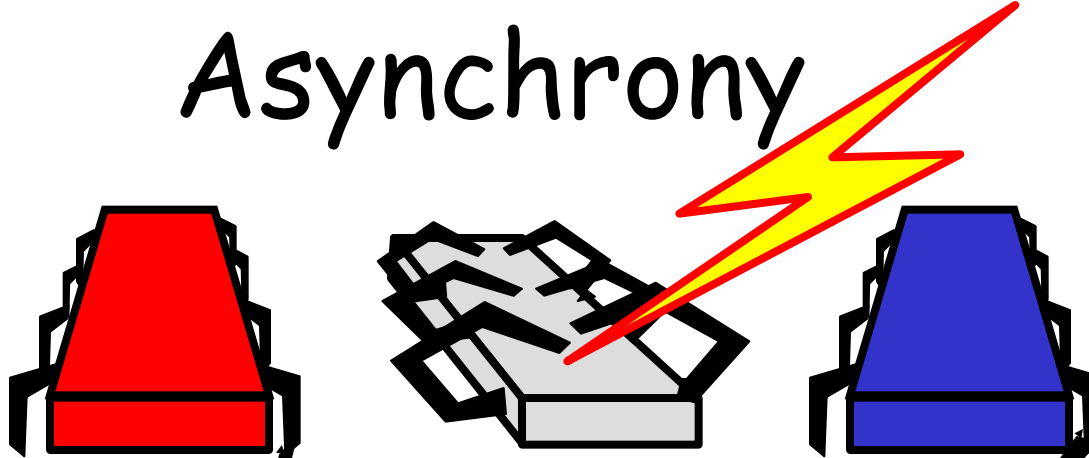
Sequential Computation



Concurrent Computation



Asynchrony



- Sudden unpredictable delays
 - Cache misses (short)
 - Page faults (long)
 - Scheduling quantum used up (really long)

Model Summary

- Multiple threads
 - Sometimes called processes
- Single shared memory
- Objects live in memory
- Unpredictable asynchronous delays

Road Map

- We are going to focus on principles first, then practice. We want to understand what we can and cannot compute before we try and write code.
 - Start with idealized models
 - Look at simplistic problems
 - Emphasize correctness over pragmatism
 - "Correctness may be theoretical, but incorrectness has practical impact"

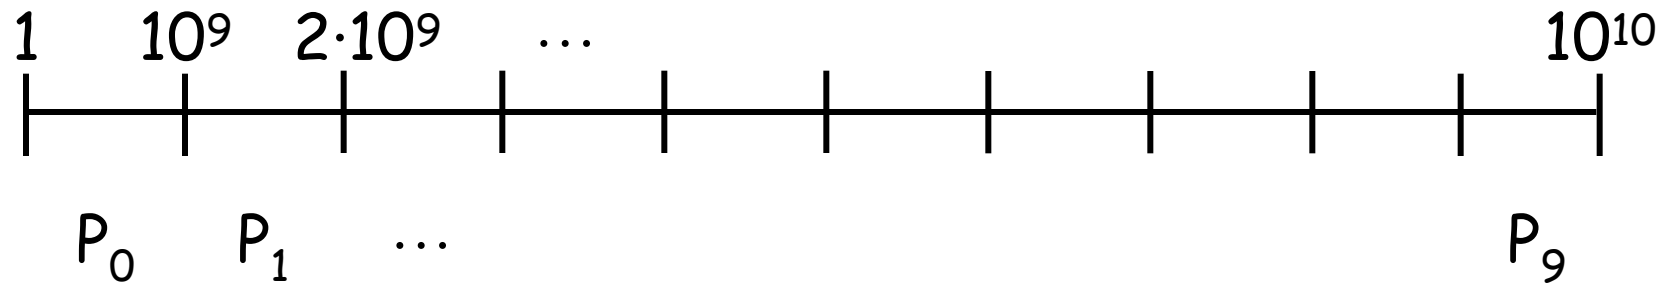
Concurrency Jargon

- Hardware
 - Processors
- Software
 - Threads, processes
- Sometimes OK to confuse them, sometimes not.

Parallel Primality Testing

- Challenge
 - Print primes from 1 to 10^{10}
- Given
 - Ten-processor multiprocessor
 - One thread per processor
- Goal
 - Get ten-fold speedup (or close)

Load Balancing



- Split the work evenly
- Each thread tests range of 10^9

Procedure for Thread i

```
void primePrint {  
    int i = ThreadID.get(); // IDs in {0..9}  
    for (j = i*109+1, j<(i+1)*109; j++) {  
        if (isPrime(j))  
            print(j);  
    }  
}
```

```
public class ThreadID {  
    private static volatile int nextID=0;  
    private static class ThreadLocalID extends ThreadLocal<Integer>{  
        protected synchronized Integer initialValue(){  
            return nextID ++;  
        }  
    }  
    private static ThreadLocalID threadID =new ThreadLocalID();  
    public static int get(){  
        return threadID.get();  
    }  
    public static void set (int index){  
        threadID.set(index);  
    }  
}
```


Issues

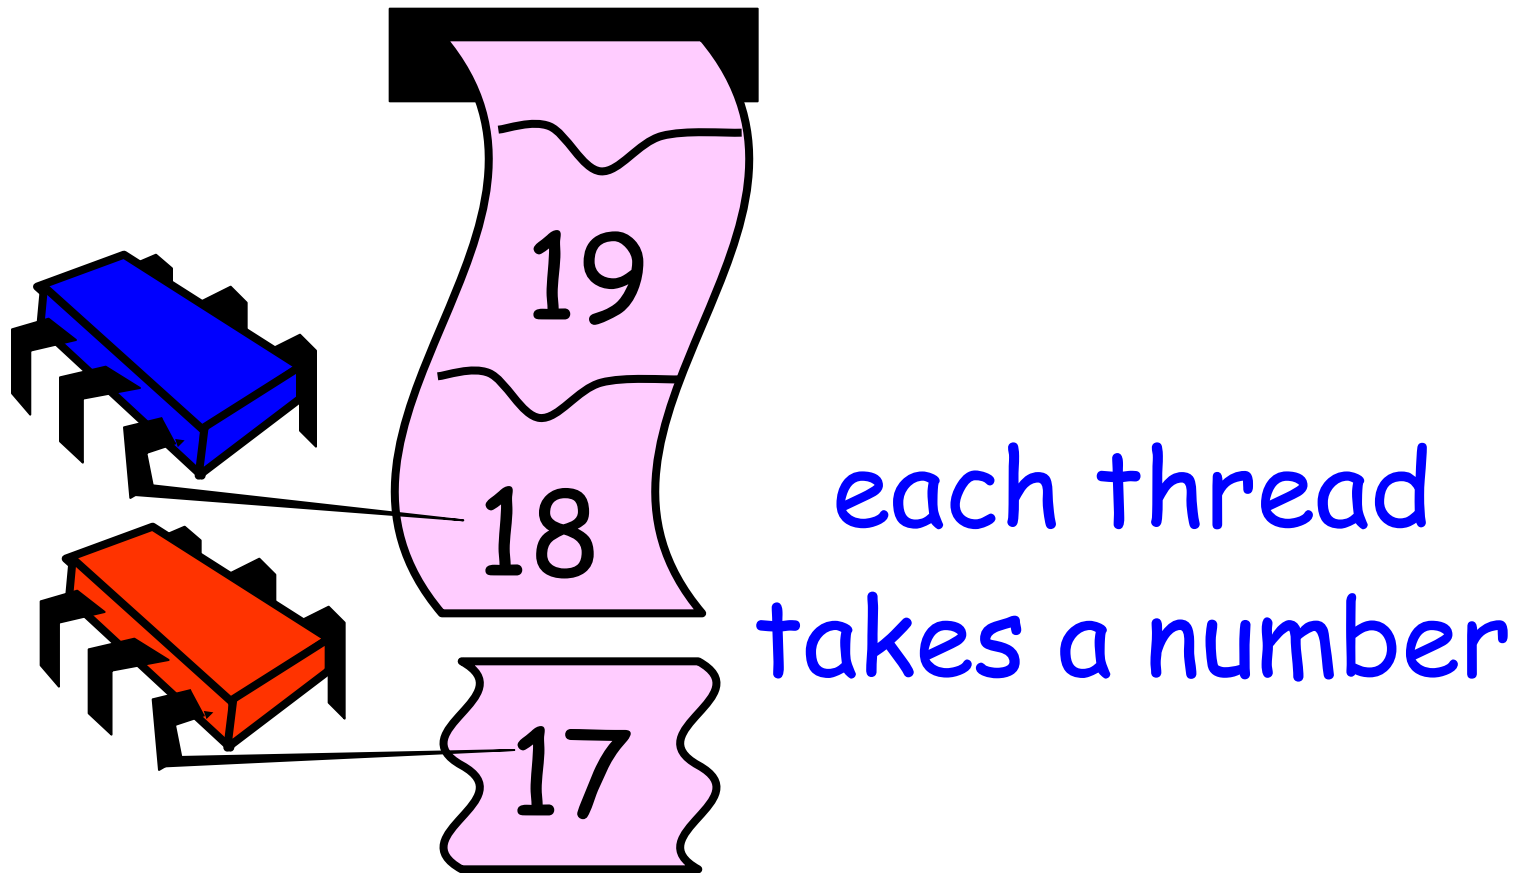
- Higher ranges have fewer primes
- Yet larger numbers harder to test
- Thread workloads
 - Uneven
 - Hard to predict

Issues

- Higher ranges have fewer primes
- Yet larger numbers harder to test
- Thread workloads
 - Uneven
 - Hard to predict
- Need dynamic load balancing



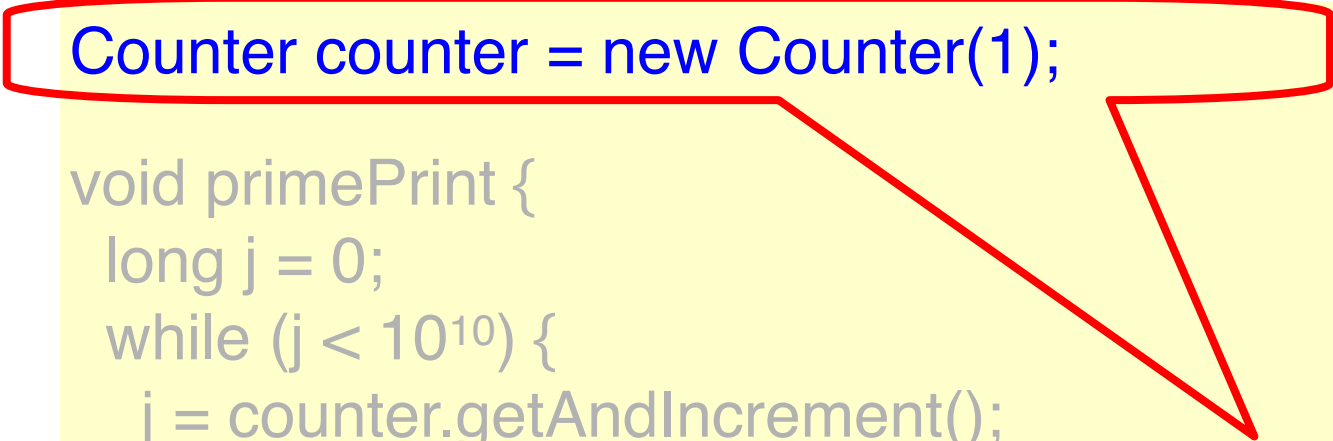
Shared Counter



Procedure for Thread i

```
counter counter = new Counter(1);  
  
void primePrint {  
    long j = 0;  
    while (j < 1010) {  
        j = counter.getAndIncrement();  
        If (j < 1010)  
            if (isPrime(j))print(j); }  
}
```

Procedure for Thread i



```
Counter counter = new Counter(1);
```

```
void primePrint {  
    long j = 0;  
    while (j < 1010) {  
        j = counter.getAndIncrement();  
        If (j < 1010)  
            if (isPrime(j)) print(j); }  
}
```

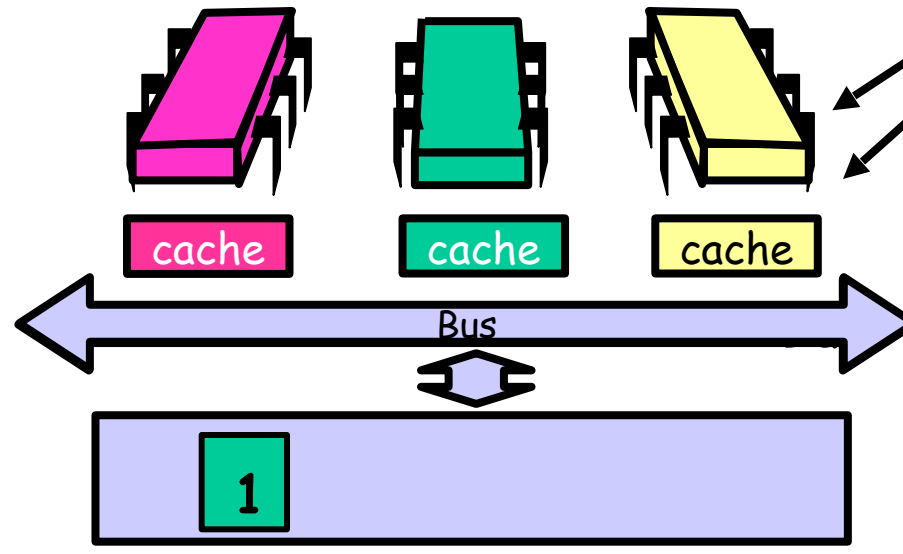
Shared counter
object

Where Things Reside

```
void primePrint {  
  int i = ThreadID.get(); // IDs in {0..9}  
  for (j = i*10^2+1, j<(i+1)*10^2; j++) {  
    if (isPrime(j))  
      print(j);  
  }  
}
```

code

Local
variables



shared
memory

shared counter

Procedure for Thread i

```
Counter counter = new Counter(1);
```

```
void primePrint {
```

```
    long j = 0;
```

```
    while (j < 1010) {
```

```
        j = counter.getAndIncrement();
```

```
        If (j < 1010)
```

```
            if (isPrime(j)) print(j);
```

```
    }
```

```
}
```

Stop when every
value taken

Procedure for Thread i

```
Counter counter = new Counter(1);
```

```
void primePrint {  
    long j = 0;  
    while (j < 1010) {
```

```
        j = counter.getAndIncrement();
```

```
        If (j < 1010)
```

```
            if (isPrime(j)) print(j);
```

```
    }
```

```
}
```

Increment & return
each new value

Counter Implementation

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        return value++;  
    }  
}
```

Counter Implementation

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        return value++;  
    }  
}
```

OK for single thread,
not for concurrent threads

What It Means

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        return value++;  
    }  
}
```

What It Means

```
public class Counter {  
    private long value;
```

```
    public long getAndIncrement() {
```

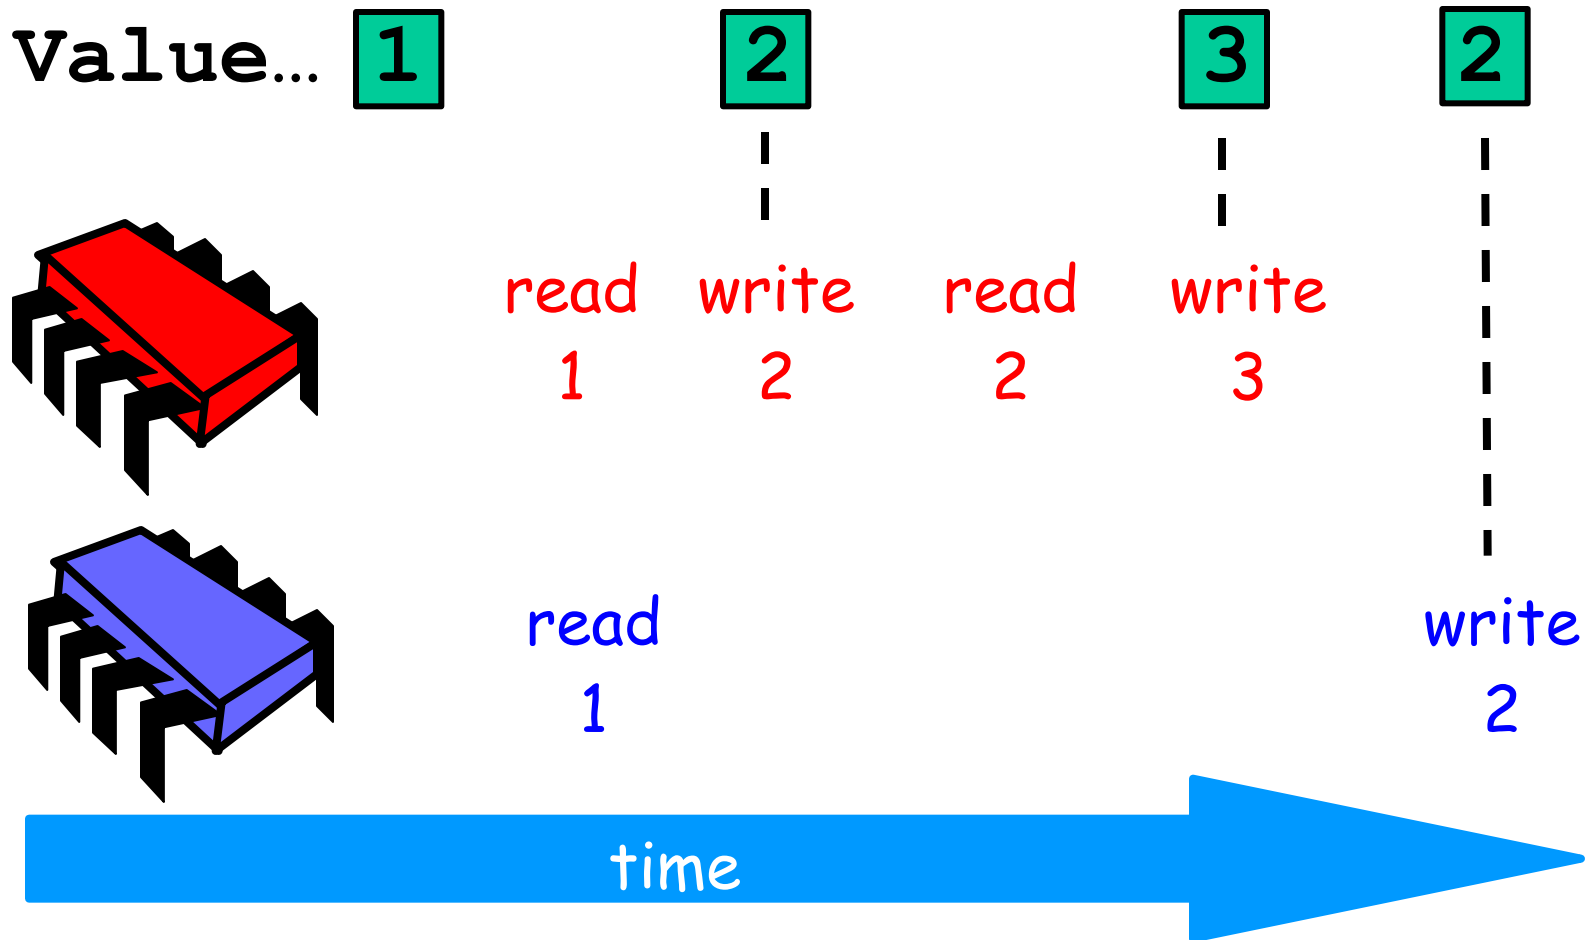
```
        return value++;
```

```
    }
```

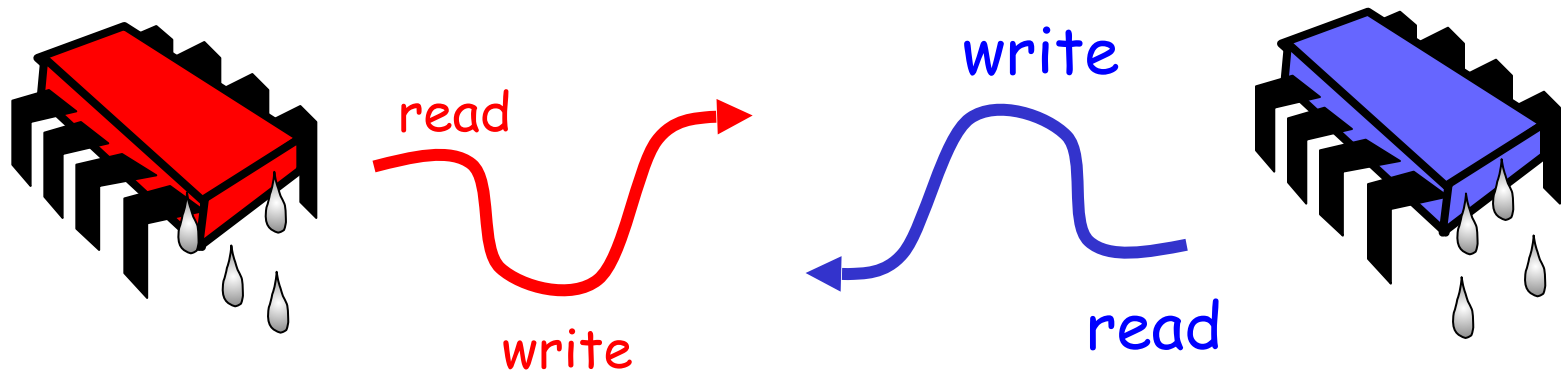
```
}
```

```
    temp = value;  
    value = temp + 1;  
    return temp;
```

Not so good...



Is this problem inherent?



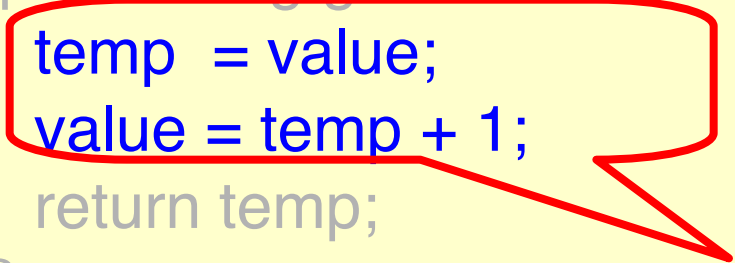
If we could only glue reads and writes...

Challenge

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        temp = value;  
        value = temp + 1;  
        return temp;  
    }  
}
```

Challenge

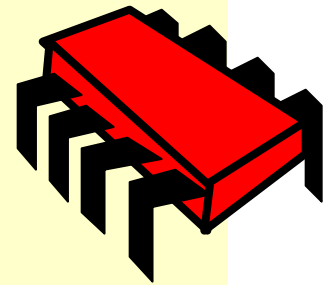
```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        temp = value;  
        value = temp + 1;  
        return temp;  
    }  
}
```



Make these steps
atomic (indivisible)

Hardware Solution

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        temp = value;  
        value = temp + 1;  
        return temp;  
    }  
}
```



ReadModifyWrite()
instruction

An Aside: Java™

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        synchronized {  
            temp = value;  
            value = temp + 1;  
        }  
        return temp;  
    }  
}
```

An Aside: Java™

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        synchronized {  
            temp = value;  
            value = temp + 1;  
        }  
        return temp;  
    }  
}
```

Synchronized block

An Aside: Java™

```
public class Counter {  
    private long value;
```

```
    public long getAndIncrement() {
```

```
        synchronized {
```

```
            temp = value;
```

```
            value = temp + 1;
```

```
        }
```

```
        return temp;
```

```
    }
```

```
}
```

Mutual Exclusion



Formalizing the Problem

- Two types of formal properties in asynchronous computation:
- Safety Properties
 - Nothing bad happens ever
- Liveness Properties
 - Something good happens eventually

Formalizing our Problem

- Mutual Exclusion
 - never simultaneously
 - This is a **safety** property
- No Deadlock
 - if only one wants in, it gets in
 - if both want in, one gets in.
 - This is a **liveness** property

Formalizing our Problem

- No starvation (many flavours)
 - if one wants in, it eventually gets in
 - if one wants in, it gets in after at most n processes enter in sc before it
- This is a **liveness** property

- Mutual Exclusion
 - Producers/Consumers
 - Readers/Writers
-
- Waiting

- Concurrent data structure
 - specification?
 - implementation non blocking, wait free

Why do we care?

- We want as much of the code as possible to execute concurrently (in parallel)
- A larger sequential part implies reduced performance
- **Amdahl's law:** this relation is not linear...

Amdahl's Law

$$\text{Speedup} = \frac{\text{Temps séquentiel}}{\text{Temps concurrent}}$$

...of computation given n CPUs instead of 1

Amdahl's Law

(normalized)
time 1 for
a single
processor
to complete
the code

Speedup=

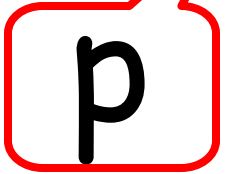
1

$$1 - p + \frac{p}{n}$$

Amdahl's Law

Speedup =
$$\frac{1}{1 - p + \frac{p}{n}}$$

Parallel fraction



Amdahl's Law

Sequential
fraction

Speedup =

Parallel
fraction

$$\frac{1}{1 - p + \frac{p}{n}}$$

Amdahl's Law

Sequential
fraction

Parallel
fraction

Speedup =

$$\frac{1}{(1 - p) + \frac{p}{n}}$$

The diagram shows the formula for Amdahl's Law: $\text{Speedup} = \frac{1}{(1 - p) + \frac{p}{n}}$. Red annotations highlight the components: a box around '1 - p' is connected by a red line to the text 'Sequential fraction'; a box around 'p' is connected by a red line to the text 'Parallel fraction'; and a box around 'n' is connected by a red line to the text 'Number of processors'.

Number of
processors

Example

- Ten processors
- 60% concurrent, 40% sequential
- How close to 10-fold speedup?

Example

- Ten processors
- 60% concurrent, 40% sequential
- How close to 10-fold speedup?

$$\text{Speedup}=2.17= \frac{1}{1 - 0.6 + \frac{0.6}{10}}$$

Example

- Ten processors
- 80% concurrent, 20% sequential
- How close to 10-fold speedup?

Example

- Ten processors
- 80% concurrent, 20% sequential
- How close to 10-fold speedup?

$$\text{Speedup}=3.57=\frac{1}{1-0.8+\frac{0.8}{10}}$$

Example

- Ten processors
- 90% concurrent, 10% sequential
- How close to 10-fold speedup?

Example

- Ten processors
- 90% concurrent, 10% sequential
- How close to 10-fold speedup?

$$\text{Speedup}=5.26= \frac{1}{1 - 0.9 + \frac{0.9}{10}}$$

Example

- Ten processors
- 99% concurrent, 01% sequential
- How close to 10-fold speedup?

Example

- Ten processors
- 99% concurrent, 01% sequential
- How close to 10-fold speedup?

$$\text{Speedup}=9.17= \frac{1}{1 - 0.99 + \frac{0.99}{10}}$$

The Moral

- Making good use of our multiple processors (cores) means
- Finding ways to effectively parallelize our code
 - Minimize sequential parts
 - Reduce idle time in which threads wait without

Multicore Programming

- This is what this course is about...
 - The % that is not easy to make concurrent yet may have a large impact on overall speedup

Plan du cours

- » Une partie théorique et une partie pratique
- » Java - Threads
- » Exclusion mutuelle
- » Objets concurrents- correction et progression
- » Java - `java.util.concurrent.atomic`
- » Objets universels
- » Implementation-Performance

Why is Concurrent Programming so Hard?

- Try preparing a seven-course banquet
 - By yourself
 - With one friend
 - With twenty-seven friends ...
- Before we can talk about programs
 - Need a language
 - Describing time and concurrency

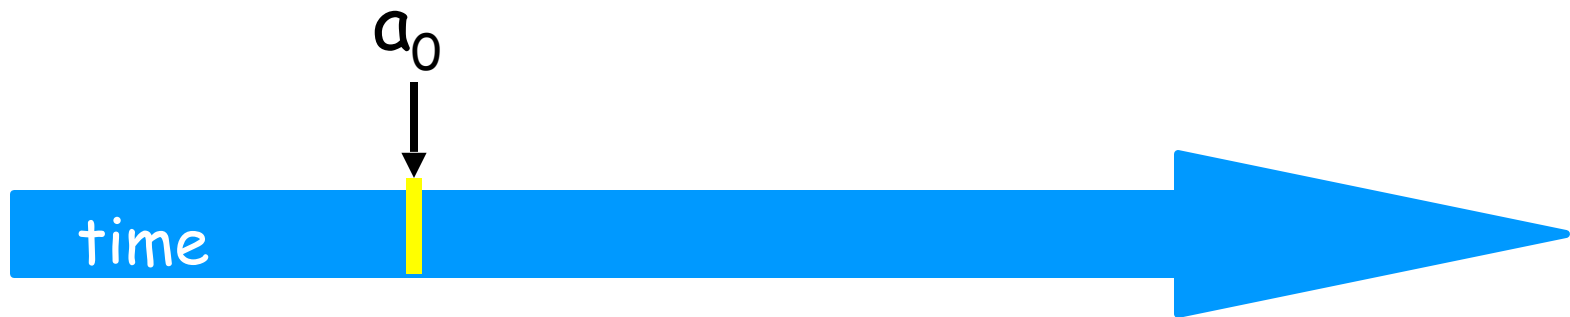
Time

- "Absolute, true and mathematical time, of itself and from its own nature, flows equably without relation to anything external." (I. Newton, 1689)
- "Time is, like, Nature's way of making sure that everything doesn't happen all at once." (Anonymous, circa 1968)



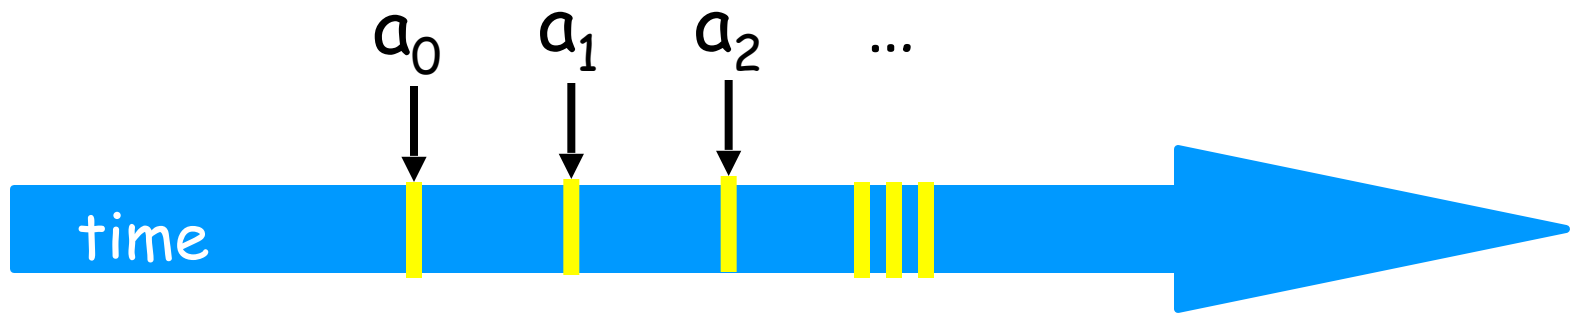
Events

- An **event** a_0 of thread A is
 - Instantaneous
 - No simultaneous events (break ties)



Threads

- A **thread** A is (formally) a sequence a_0, a_1, \dots of events
 - "Trace" model
 - Notation: $a_0 \rightarrow a_1$ indicates order

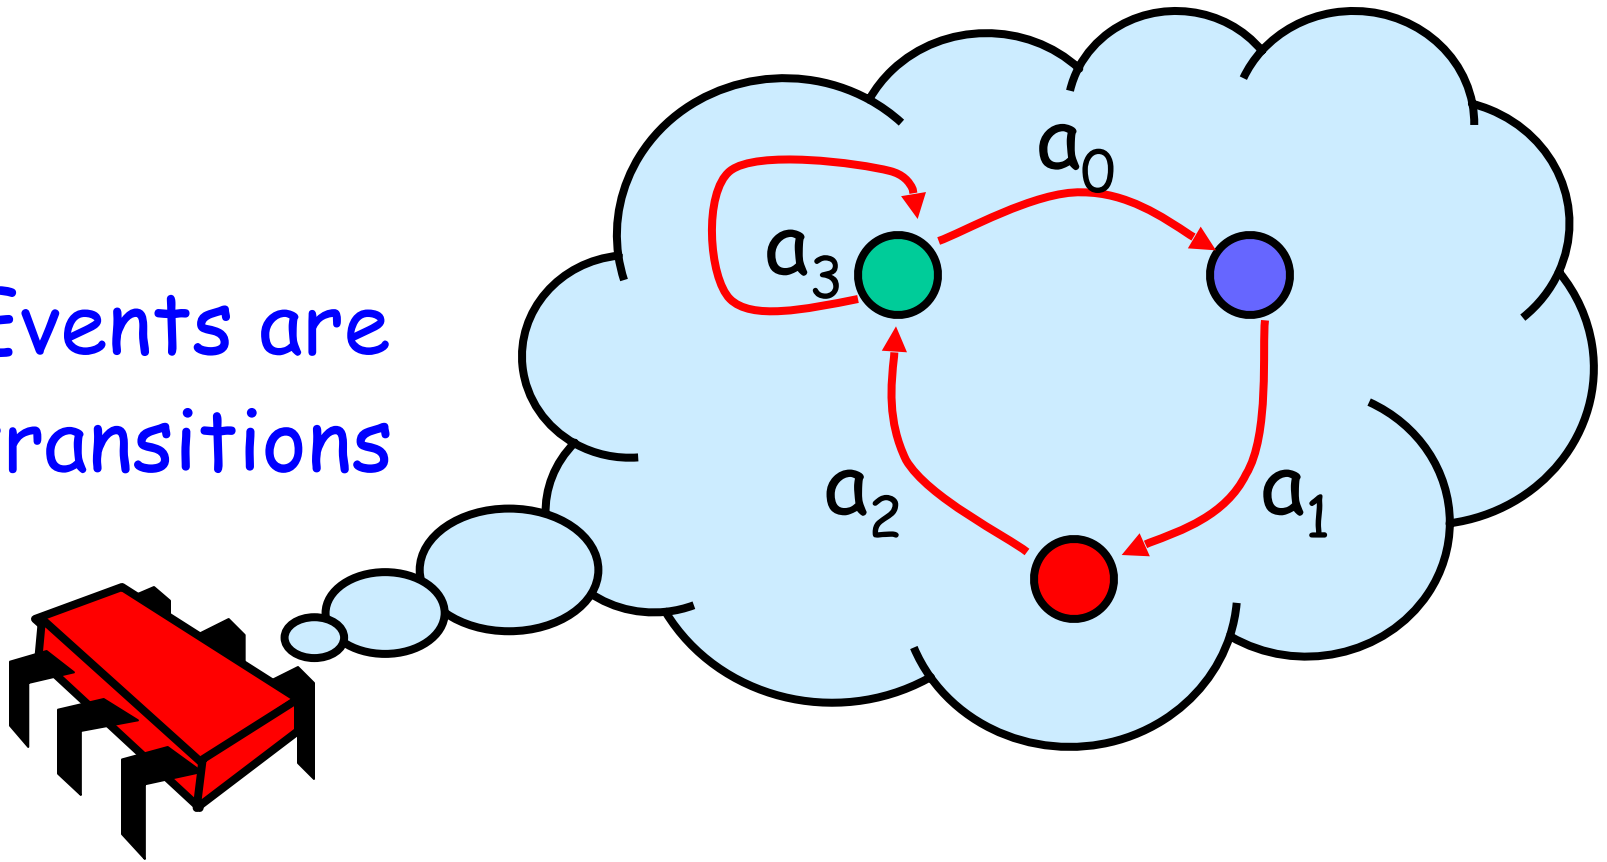


Example Thread Events

- Assign to shared variable
- Assign to local variable
- Invoke method
- Return from method
- Lots of other things ...

Threads are State Machines

Events are
transitions



States

- Thread State
 - Program counter
 - Local variables
- System state
 - Object fields (shared variables)
 - Union of thread states

Concurrency

- Thread A

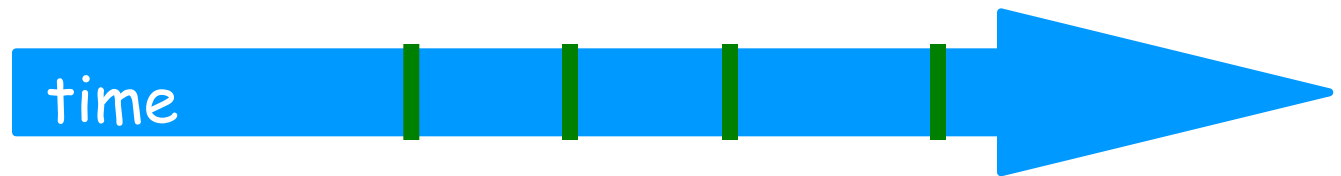


Concurrency

- Thread A



- Thread B



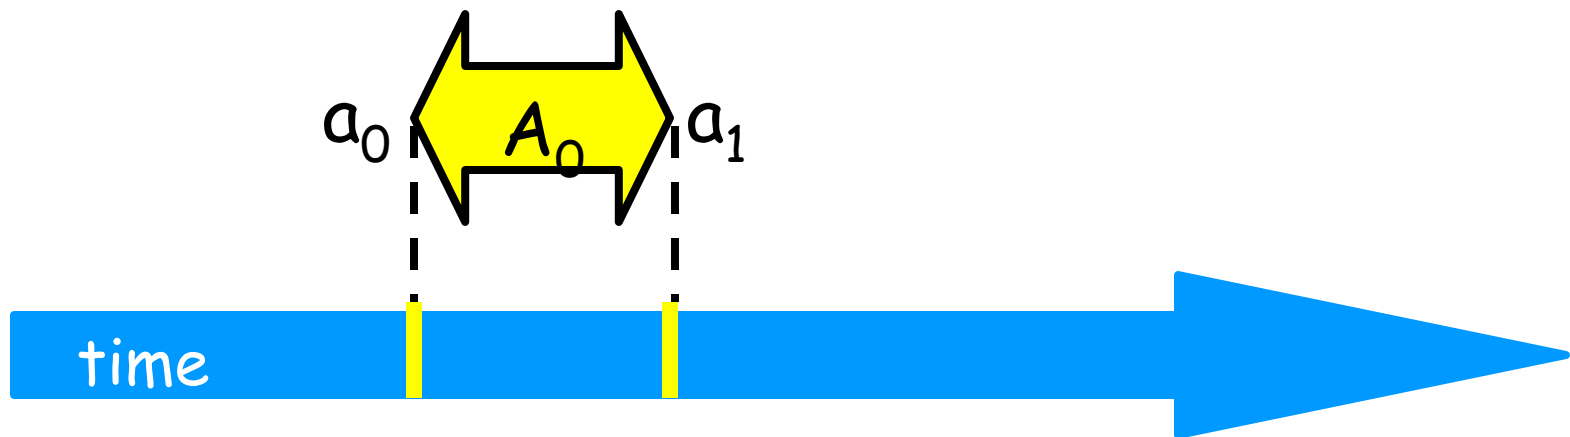
Interleavings

- Events of two or more threads
 - Interleaved
 - Not necessarily independent (why?)

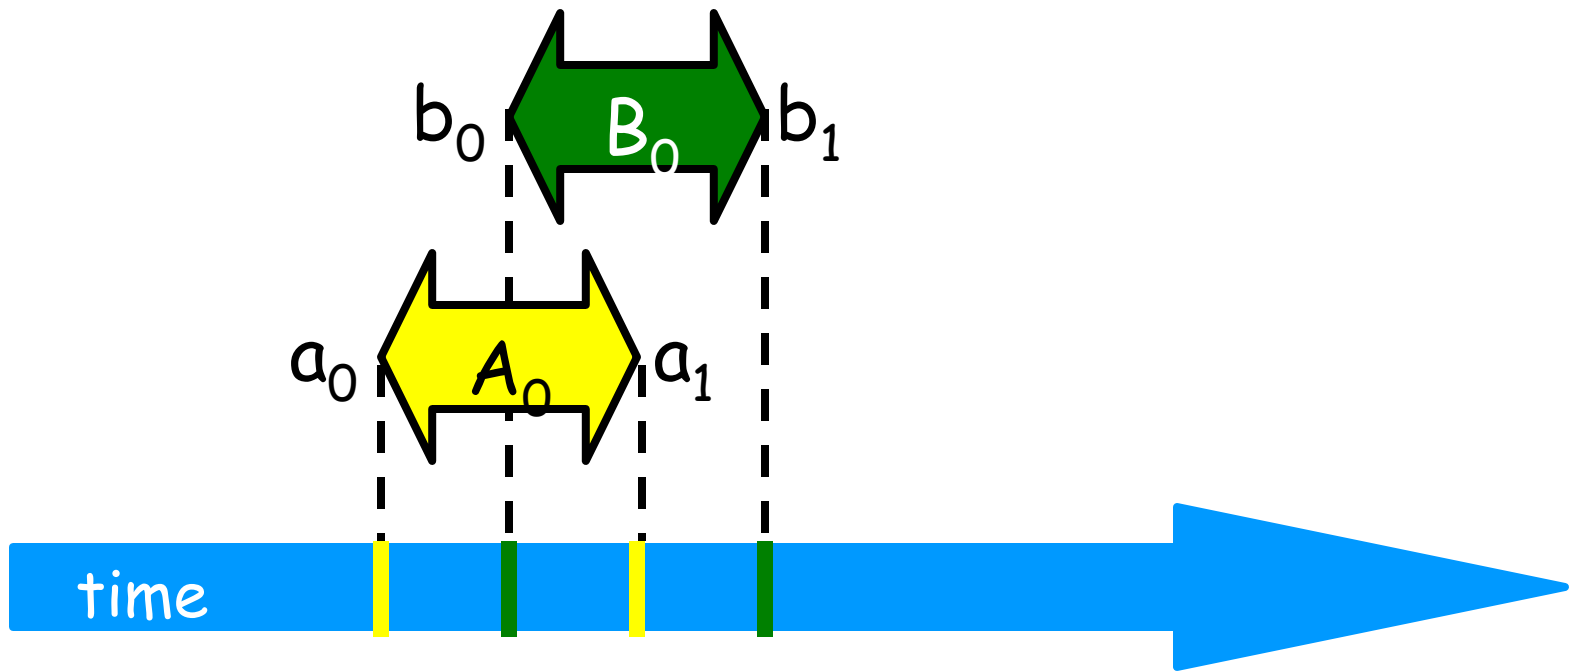


Intervals

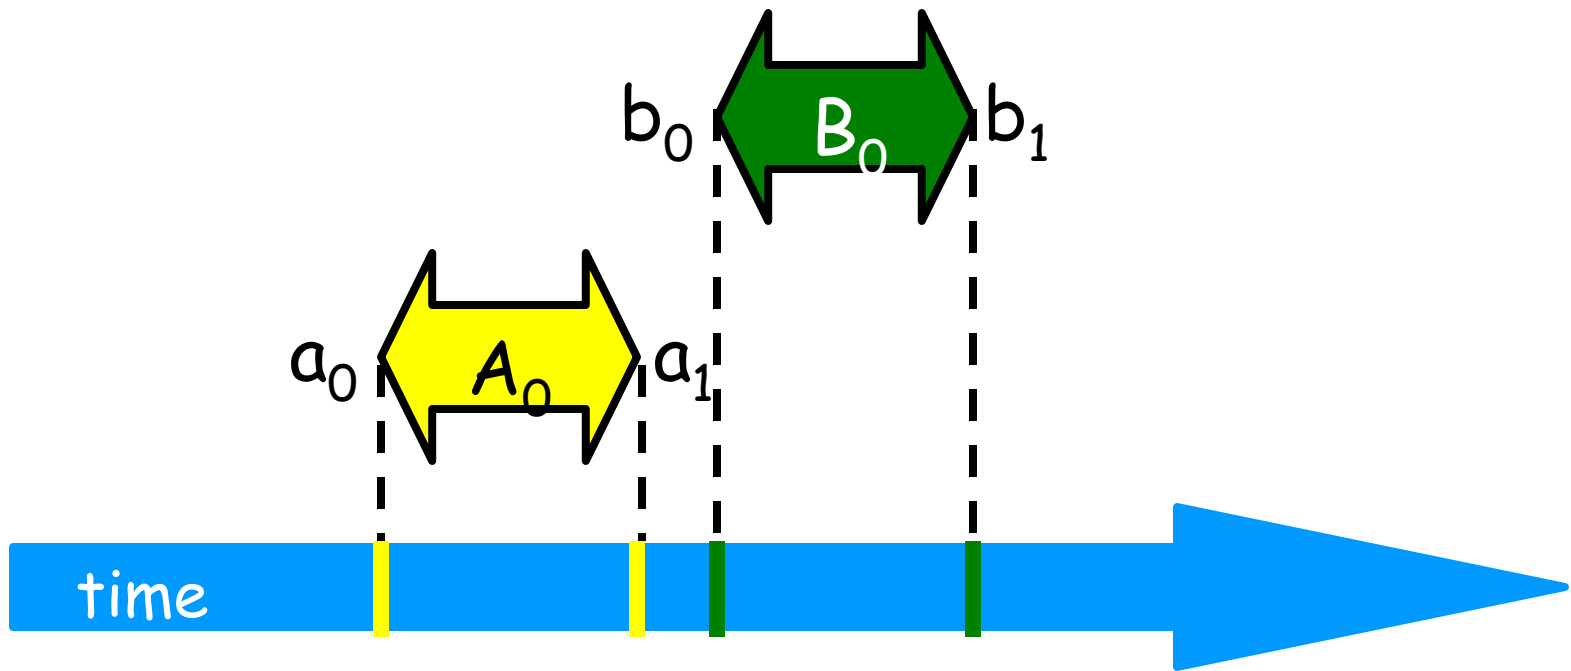
- An interval $A_0 = (a_0, a_1)$ is
 - Time between events a_0 and a_1



Intervals may Overlap

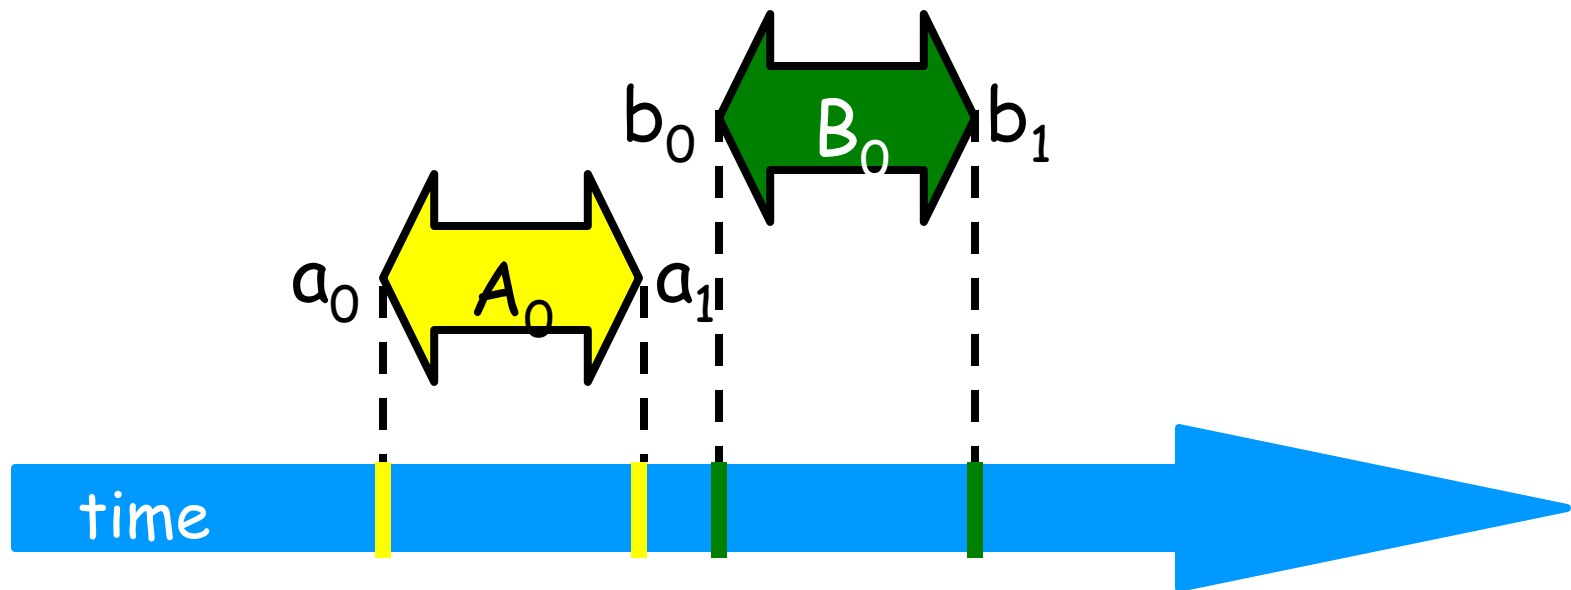


Intervals may be Disjoint

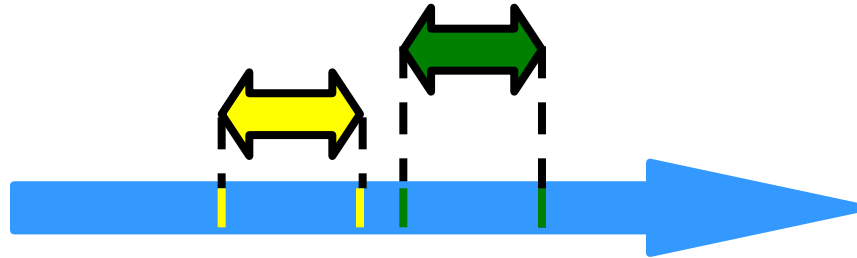


Precedence

Interval A_0 precedes interval B_0

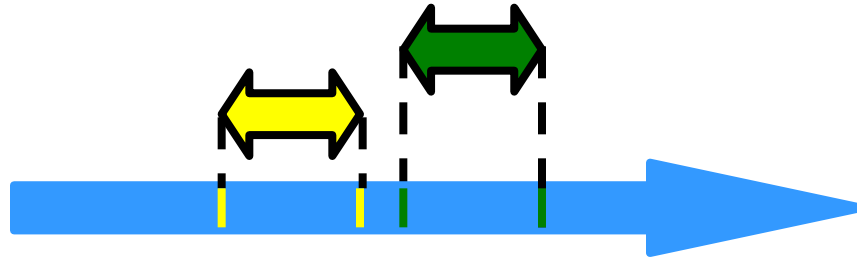


Precedence



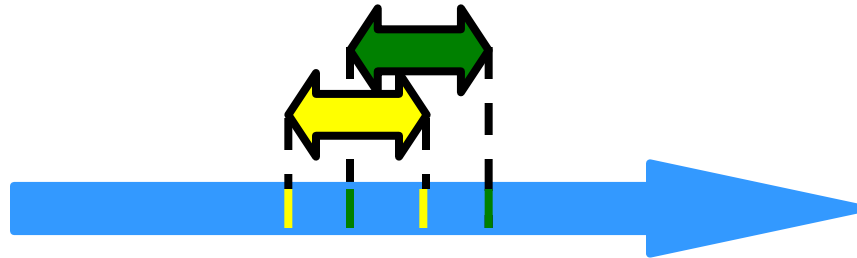
- Notation: $A_0 \rightarrow B_0$
- Formally,
 - End event of A_0 before start event of B_0
 - Also called "happens before" or "precedes"

Precedence Ordering



- Remark: $A_0 \rightarrow B_0$ is just like saying
 - 1066 AD \rightarrow 1492 AD,
 - Middle Ages \rightarrow Renaissance,
- Oh wait,
 - what about this week **vs** this month?

Precedence Ordering



- Never true that $A \rightarrow A$
- If $A \rightarrow B$ then not true that $B \rightarrow A$
- If $A \rightarrow B$ & $B \rightarrow C$ then $A \rightarrow C$
- Funny thing: $A \rightarrow B$ & $B \rightarrow A$ might both be false!

Partial Orders

(you may know this already)

- Irreflexive:
 - Never true that $A \rightarrow A$
- Antisymmetric:
 - If $A \rightarrow B$ then not true that $B \rightarrow A$
- Transitive:
 - If $A \rightarrow B$ & $B \rightarrow C$ then $A \rightarrow C$

Total Orders

(you may know this already)

- Also
 - Irreflexive
 - Antisymmetric
 - Transitive
- Except that for every distinct A, B ,
 - Either $A \rightarrow B$ or $B \rightarrow A$

Repeated Events

```
while (mumble) {  
    a0; a1;  
}
```

k-th occurrence
of event a_0

a_0^k

k-th occurrence of
interval $A_0 = (a_0, a_1)$

A_0^k

Implementing a Counter

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        temp = value;  
        value = temp + 1;  
        return temp;  
    }  
}
```

Make these steps
indivisible using
locks

Locks (Mutual Exclusion)

```
public interface Lock {  
  
    public void lock();  
  
    public void unlock();  
}
```


Locks (Mutual Exclusion)

```
public interface Lock {
```

```
    public void lock();
```

acquire lock

```
    public void unlock();  
}
```

Locks (Mutual Exclusion)

```
public interface Lock {
```

```
    public void lock();
```

acquire lock

```
    public void unlock();
```

release lock

```
}
```

Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```

Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```

acquire Lock

Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```

Release lock
(no matter what)


Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```

Critical
section

Mutual Exclusion







Mutual Exclusion

- Let CS_i^k , , be thread i 's k -th critical section execution





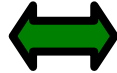

Mutual Exclusion

- Let CS_i^k, \Leftrightarrow , be thread i 's k -th critical section execution
- And CS_j^m, \Leftrightarrow , be thread j 's m -th critical section execution

Mutual Exclusion

- Let CS_i^k , , be thread i's k-th critical section execution
- And CS_j^m , , be j's m-th execution
- Then either
 -   or  



Mutual Exclusion

- Let CS_i^k  be thread i 's k -th critical section execution
- And CS_j^m  be j 's m -th execution
- Then either
 -   or  



$CS_i^k \rightarrow CS_j^m$

Mutual Exclusion

- Let CS_i^k  be thread i 's k -th critical section execution
- And CS_j^m  be j 's m -th execution
- Then either

-   or  

$CS_i^k \rightarrow CS_j^m$

$CS_j^m \rightarrow CS_i^k$

Deadlock-Free



- If some thread calls **lock()**
 - And never returns
 - Then other threads must complete **lock()** and **unlock()** calls infinitely often
- System as a whole makes progress
 - Even if individuals starve

Starvation-Free



- If some thread calls `lock()`
 - It will eventually return
- Individual threads make progress

Two-Thread vs n -Thread Solutions

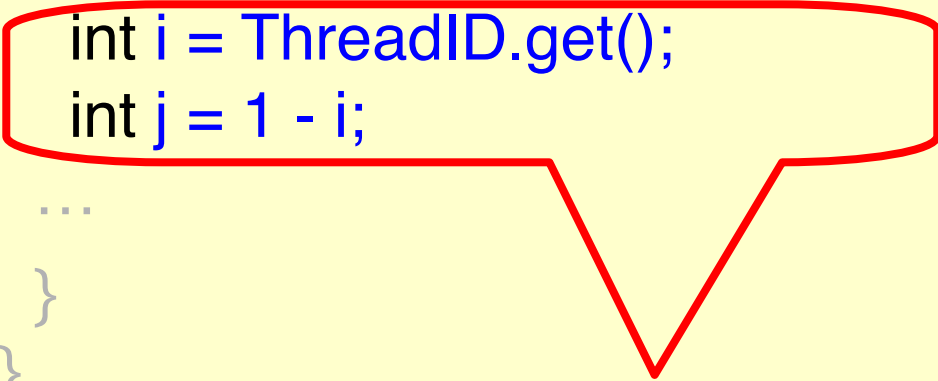
- Two-thread solutions first
 - Illustrate most basic ideas
 - Fits on one slide
- Then n-Thread solutions

Two-Thread Conventions

```
class ... implements Lock {  
    ...  
    // thread-local index, 0 or 1  
    public void lock() {  
        int i = ThreadID.get();  
        int j = 1 - i;  
        ...  
    }  
}
```


Two-Thread Conventions

```
class ... implements Lock {  
    ...  
    // thread-local index, 0 or 1  
    public void lock() {  
        int i = ThreadID.get();  
        int j = 1 - i;  
        ...  
    }  
}
```



Henceforth: *i* is current
thread, *j* is other thread

LockOne

```
class LockOne implements Lock {  
    private boolean[] flag =  
        new boolean[2];  
  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
  
    public void unlock(){  
        flag[i]= false;  
    }  
}
```

JAVA like

LockOne

```
class LockOne implements Lock {  
    private boolean[] flag =  
        new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

Set my flag

LockOne

```
class LockOne implements Lock {  
    private boolean[] flag =  
        new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

Set my flag

Wait for other
flag to go false

LockOne Satisfies Mutual Exclusion

- Assume CS_A^j overlaps CS_B^k
- Consider each thread's last (j-th and k-th) read and write in the lock() method before entering
- Derive a contradiction

From the Code

```
class LockOne implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

- $\text{write}_A(\text{flag}[A]=\text{true}) \rightarrow$
 $\text{read}_A(\text{flag}[B] \neq \text{false}) \rightarrow \text{CS}_A$
- $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow$
 $\text{read}_B(\text{flag}[A] \neq \text{false}) \rightarrow \text{CS}_B$

```
class LockOne implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow$
 $\text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow$
 $\text{write}_A(\text{flag}[B] = \text{true})$

Combining

```
class LockOne implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

- Assumptions:

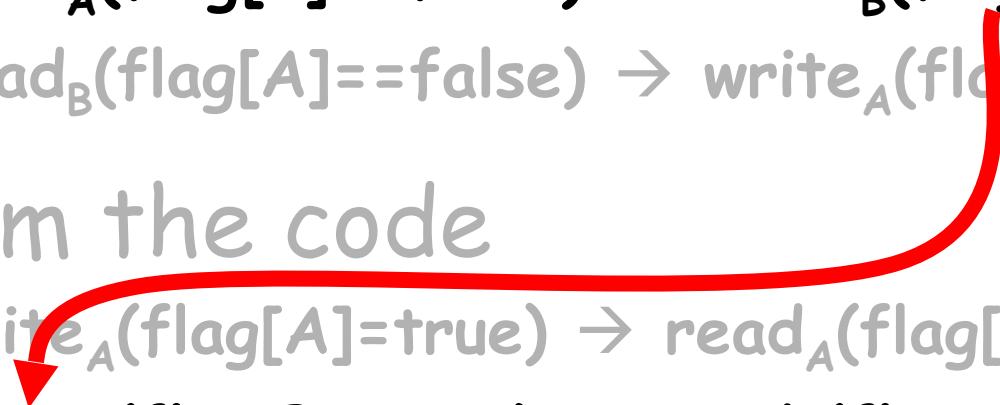
- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

Combining

```
class LockOne implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

- Assumptions:
 - $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
 - $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$
 - From the code
 - $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
 - $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$
- 

Combining

- Assumptions:

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

Combining

- Assumptions:

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

Combining

- Assumptions:

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$

- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

Combining

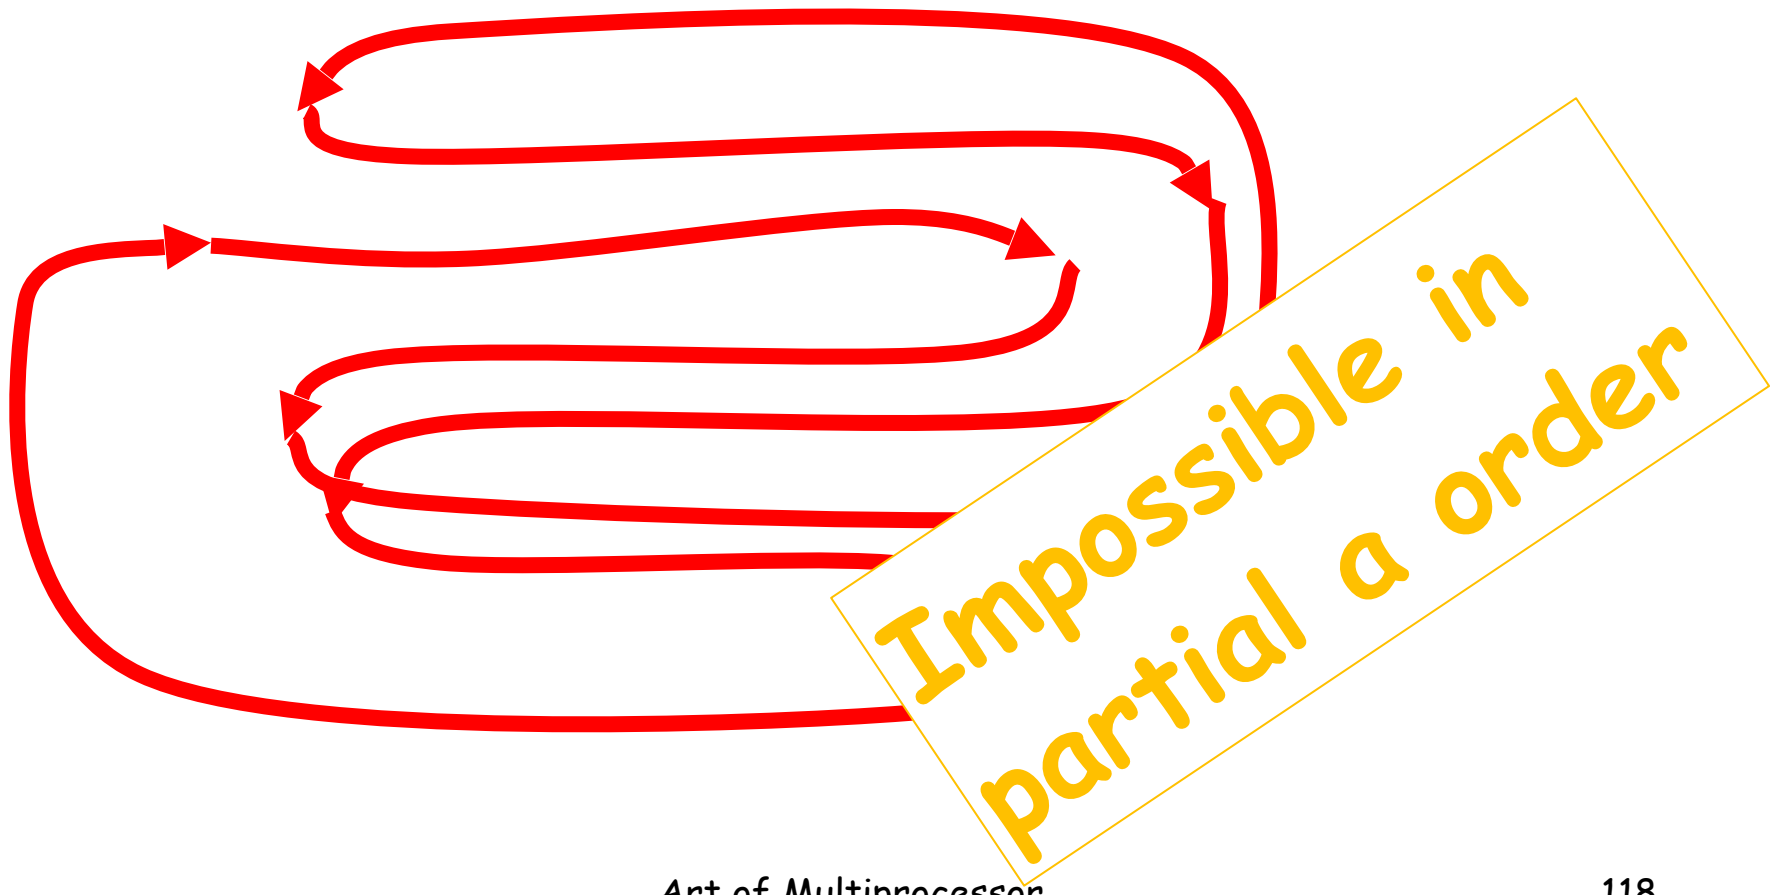
- Assumptions:

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

Cycle!



Deadlock Freedom

- LockOne Fails deadlock-freedom
 - Concurrent execution can deadlock

```
flag[i] = true;   flag[j] = true;  
while (flag[j]){ while (flag[i]){
```

- Sequential executions OK

LockTwo

```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {}  
    }  
  
    public void unlock() {}  
}
```


LockTwo

```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {}  
    }  
  
    public void unlock() {}  
}
```

**Let other go
first**

LockTwo

```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {};  
    }  
  
    public void unlock() {}  
}
```

**Wait for
permission**

LockTwo

```
public class Lock2 implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {}  
    }
```

Nothing to do

```
    public void unlock() {}  
}
```

LockTwo Claims

- Satisfies mutual exclusion
 - If thread i in CS
 - Then $\text{victim} == j$
 - Cannot be both 0 and 1
- Not deadlock free
 - Sequential execution deadlocks
 - Concurrent execution does not

```
public void LockTwo() {  
    victim = i;  
    while (victim == i) {}  
}
```

Peterson's Algorithm

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {}  
}  
public void unlock() {  
    flag[i] = false;  
}
```

Peterson's Algorithm

Announce I'm
interested

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {}  
}  
public void unlock() {  
    flag[i] = false;  
}
```

Peterson's Algorithm

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {}  
}  
public void unlock() {  
    flag[i] = false;  
}
```

Announce I'm
interested

Defer to other

Peterson's Algorithm

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {}  
}  
public void unlock() {  
    flag[i] = false;  
}
```

Announce I'm
interested
Defer to other

Wait while other
interested & I'm
the victim

Peterson's Algorithm

```
public void lock() {
```

```
    flag[i] = true;
```

```
    victim = i;
```

```
    while (flag[j] && victim == i) {};
```

```
}
```

```
public void unlock() {
```

```
    flag[i] = false;
```

```
}
```

Announce I'm
interested
Defer to other

Wait while other
interested & I'm
the victim

No longer
interested

Mutual Exclusion

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {}  
}
```

- If thread **0** in critical section,
 - flag[0]=true,
 - victim = 1
- If thread **1** in critical section,
 - flag[1]=true,
 - victim = 0

Cannot both be true

From the Code

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {}  
}
```

- $\text{write}_A(\text{flag}[A]=\text{true}) \rightarrow$

$\text{write}_A(\text{victim}=A) \rightarrow$

$\text{read}_A(\text{flag}[B]) \rightarrow$

$\text{read}_A(\text{victim}) \rightarrow CS_A$

- Idem for B

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {}  
}
```

- (A was the last thread to write to the victim field)
- $\text{write}_B(\text{victim}=B) \rightarrow \text{write}_A(\text{victim}=A)$

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {}  
}
```

- $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow (\text{code})$
 $\text{write}_B(\text{victim}=B) \rightarrow (\text{hyp ordre})$
 $\text{write}_A(\text{victim}=A) \rightarrow (A \text{ en CS})$
 $\text{read}_A(\text{flag}[B]=\text{false})$
- Contradict: no other write to Flag[B] was performs
before the Cs exec

Starvation Free

- Thread *i* blocked only if *j* repeatedly re-enters so that

flag[j] == true and victim == i

- When *j* re-enters
 - it sets victim to *j*.
 - So *i* gets in

```
public void lock() {  
    flag[i] = true;  
    victim  = i;  
    while (flag[j] && victim == i) {}  
}  
  
public void unlock() {  
    flag[i] = false;  
}
```

Deadlock Free

```
public void lock() {  
    ...  
    while (flag[j] && victim == i) {};
```

- Thread blocked
 - only at while loop
 - only if it is the victim
- One or the other must not be the victim

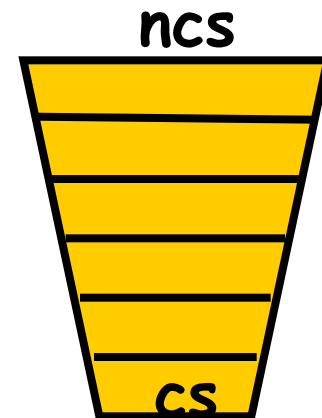
```
public void lock() {  
    ...  
    while (flag[j] && victim == i) {};
```

**Pas d'atomicité du &&
évaluation des 2 variables indépendamment**

The Filter Algorithm for n Threads

There are $n-1$ "waiting rooms" called levels

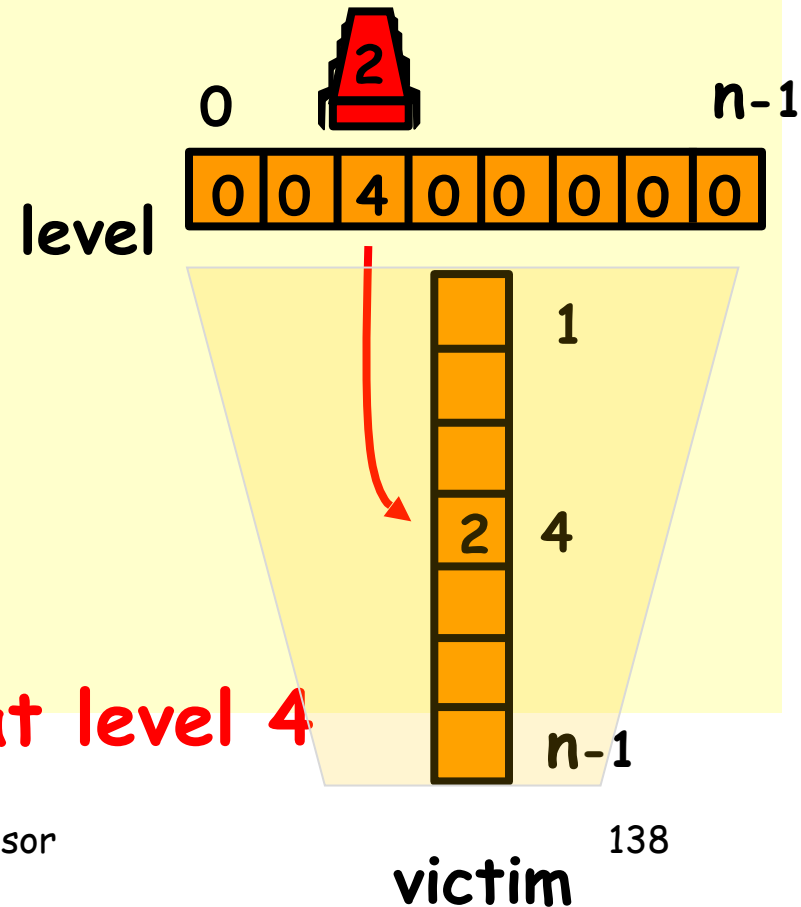
- At each level
 - At least one enters level
 - At least one blocked if many try
- Only one thread makes it through



Filter

```
class Filter implements Lock {  
    int[] level; // level[i] for thread i  
    int[] victim; // victim[L] for level L
```

```
    public Filter(int n) {  
        level = new int[n];  
        victim = new int[n];  
        for (int i = 0; i < n; i++) {  
            level[i] = 0;  
        }  
        ...  
    }
```



Thread 2 at level 4

Filter

```
class Filter implements Lock {  
    ...  
  
    public void lock(){  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists$  k != i level[k] >= L) && victim[L] == i){};  
        }  
    }  
    public void unlock() {  
        level[i] = 0;  
    }  
}
```

Filter

```
class Filter implements Lock {  
    ...  
  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists$  k != i) level[k] >= L) &&  
                victim[L] == i){};  
        }  
    }  
    public void release(int i) {  
        level[i] = 0;  
    }  
}
```

One level at a time

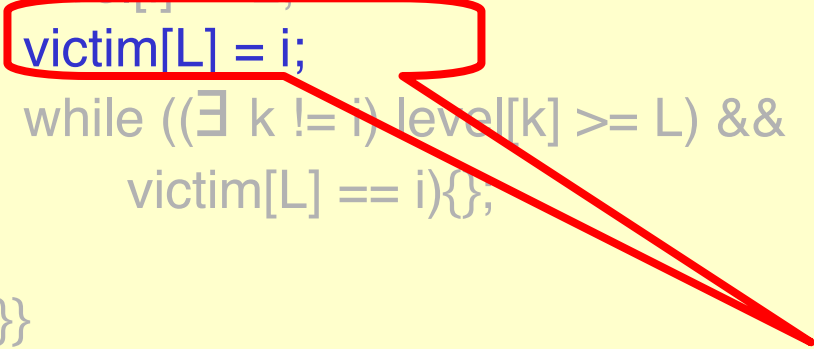
Filter

```
class Filter implements Lock {  
    ...  
  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists$  k != i) level[k] >= L) &&  
                victim[L] == i){}; // busy wait  
        }  
    }  
    public void release(int i) {  
        level[i] = 0;  
    }  
}
```

Announce
intention to enter
level L

Filter

```
class Filter implements Lock {  
    int level[n];  
    int victim[n];  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists$  k != i) level[k] >= L) &&  
                victim[L] == i){},  
        }  
    }  
    public void release(int i) {  
        level[i] = 0;  
    }  
}
```



*Give priority to
anyone but me*

Filter

Wait as long as someone else is at same or higher level, and I'm designated victim

```
public void lock() {  
    for (int L = 1; L < n; L++) {  
        level[i] = L;  
        victim[L] = i;  
        while (( $\exists$  k != i) level[k] >= L) &&  
            victim[L] == i){};  
    }  
}  
public void release(int i) {  
    level[i] = 0;  
}
```

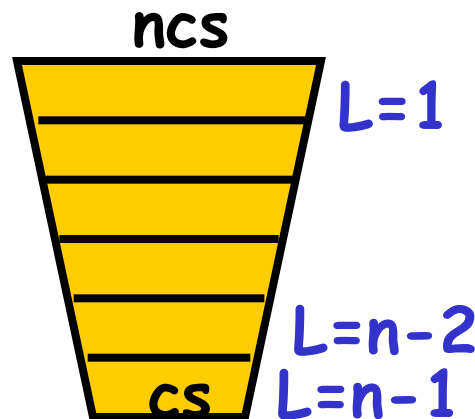
Filter

```
class Filter implements Lock {  
    int level[n];  
    int victim[n];  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists$  k != i) level[k] >= L) &&  
                victim[L] == i);  
        }  
    }
```

Thread enters level L when it completes
the loop

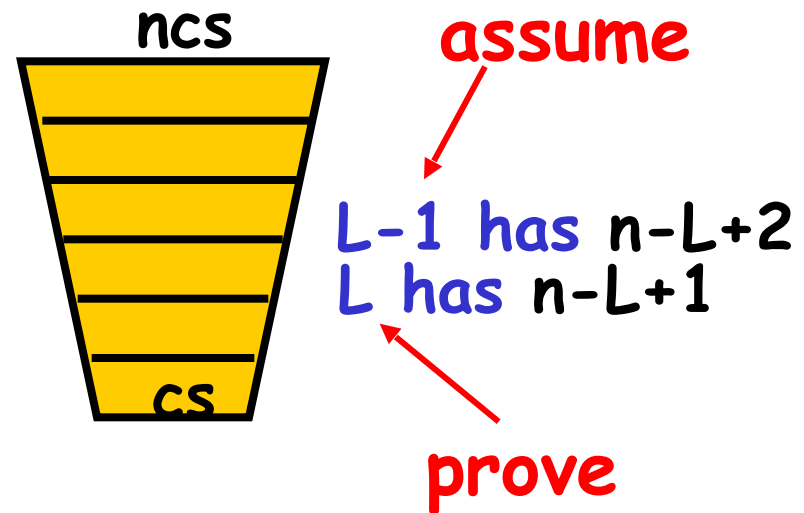
Claim

- Start at level $L=1$
- At most $n-L+1$ threads enter level L
- Mutual exclusion after level $L=n-1$

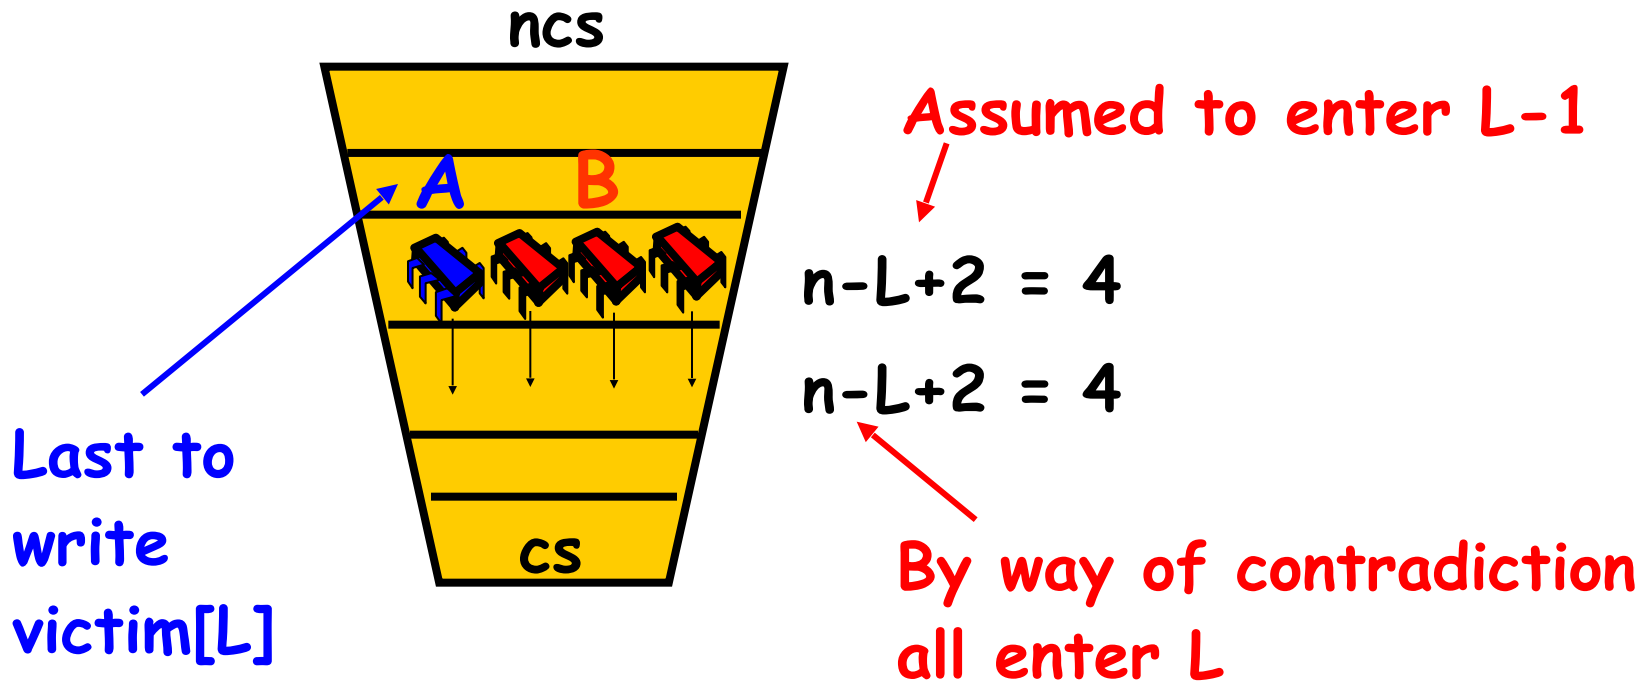


Induction Hypothesis

- No more than $n-L+1$ at level L
- Induction step: by contradiction
- By ind: no more than $n-(L-1)+1$ threads at level $L-1$
- A last to write `victim[L]`
- B is any other thread at level L



Proof Structure



Show that A must have seen B in level[L] and since victim[L] == A could not have entered

From the Code

(1) $\text{write}_B(\text{level}[B]=L) \rightarrow \text{write}_B(\text{victim}[L]=B)$

```
public void lock() {  
    for (int L = 0, L < n, L++) {  
        level[i] = L;  
        victim[L] = i;  
        while (( $\exists$  k != i) level[k] >= L)  
            && victim[L] == i) {};  
    }  
}
```

From the Code

(2) $\text{write}_A(\text{victim}[L]=A) \rightarrow \text{read}_A(\text{level}[B])$

```
public void lock() {  
    for (int L = 0; L < n; L++) {  
        level[i] = L;  
        victim[L] = i;  
        while (( $\exists k \neq i$ ) level[k] >= L)  
            && victim[L] == i) {}  
    }  
}
```

By Assumption

(3) $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$

By assumption, *A* is the last thread to write **victim[L]**

```
public void lock() {  
    for (int L = 0; L < n; L++) {  
        level[i] = L;  
        victim[L] = i;  
        while (( $\exists k \neq i$  level[k] >= L)  
                && victim[L] == i) {}  
    }  
}
```

Combining Observations

- (1) $\text{write}_B(\text{level}[B]=L) \rightarrow \text{write}_B(\text{victim}[L]=B)$
- (3) $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$
- (2) $\text{write}_A(\text{victim}[L]=A) \rightarrow \text{read}_A(\text{level}[B])$

```
public void lock() {  
    for (int L = 0; L < n; L++) {  
        level[i] = L;  
        victim[L] = i;  
        while (( $\exists k \neq i$ ) level[k] >= L  
                && victim[L] == i) {}  
    }  
}
```

Combining Observations

- (1) $\text{write}_B(\text{level}[B]=L) \rightarrow$
- (3) $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$
- (2) $\rightarrow \text{read}_A(\text{level}[B])$

```
public void lock() {  
    for (int L = 0; L < n; L++) {  
        level[i] = L;  
        victim[L] = i;  
        while (( $\exists k \neq i$ ) level[k] >= L  
                && victim[L] == i) {}  
    }  
}
```


Combining Observations

- (1) $\text{write}_B(\text{level}[B]=L) \rightarrow$
(3) $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$
(2) $\rightarrow \text{read}_A(\text{level}[B])$

Thus, A read $\text{level}[B] \geq L$,
A was last to write $\text{victim}[L]$,
so it could not have entered level L!

No Starvation

- Filter Lock satisfies properties:
 - Just like Peterson Alg at any level
 - So no one starves
- But what about fairness?
 - Threads can be overtaken by others

Bounded Waiting

- Want stronger fairness guarantees
- Thread not “overtaken” too much
- Need to adjust definitions

Bounded Waiting

- Divide `lock()` method into 2 parts:
 - Doorway interval:
 - Written D_A
 - always finishes in finite steps
 - Waiting interval:
 - Written W_A
 - may take unbounded steps

r-Bounded Waiting

- For threads A and B :
 - If $D_A^k \rightarrow D_B^j$
 - A 's k -th doorway precedes B 's j -th doorway
 - Then $CS_A^k \rightarrow CS_B^{j+r}$
 - A 's k -th critical section precedes B 's $(j+r)$ -th critical section
 - B cannot overtake A by more than r times
- First-come-first-served means $r = 0$.

Fairness

- Filter Lock satisfies properties:
 - No one starves
 - But very weak fairness

Bakery Algorithm

- Provides First-Come-First-Served
- How?
 - Take a "number"
 - Wait until lower numbers have been served
- Lexicographic order
 - $(a,i) > (b,j)$
 - If $a > b$, or $a = b$ and $i > j$

Bakery Algorithm

```
class Bakery implements Lock {  
    boolean[] flag;  
    Label[] label;  
    public Bakery (int n) {  
        flag = new boolean[n];  
        label = new Label[n];  
        for (int i = 0; i < n; i++) {  
            flag[i] = false; label[i] = 0;  
        }  
    }  
    ...  
}
```


Bakery Algorithm

```
class Bakery implements Lock {
```

```
    boolean[] flag;
```

```
    Label[] label;
```

```
    public Bakery (int n) {
```

```
        flag = new boolean[n];
```

```
        label = new Label[n];
```

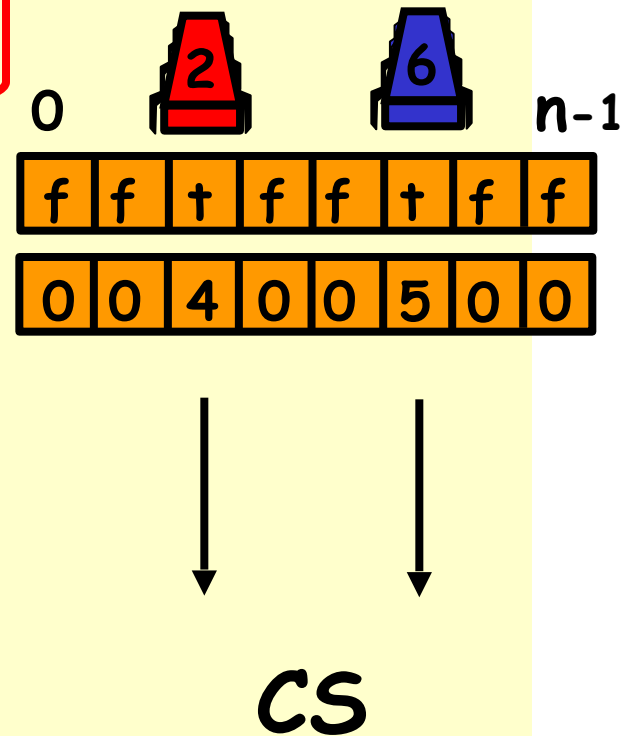
```
        for (int i = 0; i < n; i++) {
```

```
            flag[i] = false; label[i] = 0;
```

```
        }
```

```
    }
```

```
    ...
```



Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists$  k flag[k]  
                && (label[i],i) > (label[k],k));  
    }
```

Bakery Algorithm

```
class Bakery implements Lock {
```

```
...
```

```
public void lock() {
```

```
    flag[i] = true;
```

```
    label[i] = max(label[0], ..., label[n-1]) + 1;
```

```
    while (∃ k flag[k]
```

```
        && (label[i], i) > (label[k], k));
```

```
}
```

Doorway

Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
            && (label[i],i) > (label[k],k));  
    }  
}
```

I'm interested

Bakery Algorithm

Take increasing
label (read labels
in some arbitrary
order)

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
            && (label[i],i) > (label[k],k));  
    }
```

Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i],i) > (label[k],k));  
    }
```

Someone is
interested



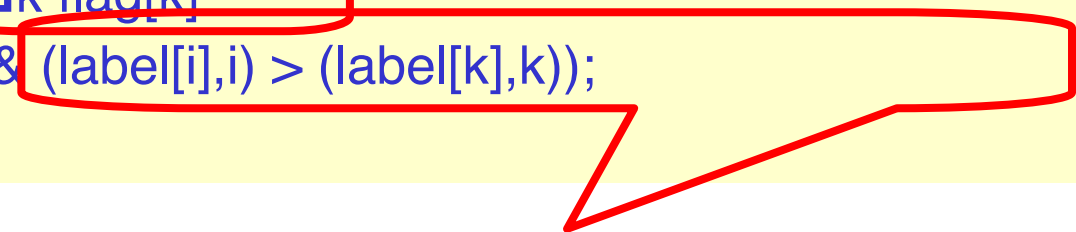
Bakery Algorithm

```
class Bakery implements Lock {  
    boolean flag[n];  
    int label[n];  
  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
            && (label[i],i) > (label[k],k));  
    }
```

**Someone is
interested**



**With lower (label,i)
in lexicographic order**



Bakery Algorithm

```
class Bakery implements Lock {
```

```
    ...
```

```
    public void unlock() {  
        flag[i] = false;  
    }  
}
```


Bakery Algorithm

```
class Bakery implements Lock {
```

```
...
```

```
public void unlock() {
```

```
    flag[i] = false;
```

```
}
```

```
}
```

**No longer
interested**



labels are always increasing

No Deadlock

- There is always one thread with earliest label
- Ties are impossible (why?)

First-Come-First-Served

- If $D_A \rightarrow D_B$ then A's label is smaller
- And:
 - $\text{write}_A(\text{label}[A]) \rightarrow$
 $\text{read}_B(\text{label}[A]) \rightarrow$
 $\text{write}_B(\text{label}[B]) \rightarrow$
 $\text{read}_B(\text{flag}[A])$
- So B is locked out while $\text{flag}[A]$ is true

```
class Bakery implements Lock {  
  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0],  
                        ..., label[n-1]) + 1;  
        while ( $\exists k$  flag[k]  
                && (label[i], i) > (label[k], k));  
    }  
}
```

Mutual Exclusion

- Suppose *A* and *B* in CS together
- Suppose *A* has earlier label
- When *B* entered, it must have seen
 - $\text{flag}[A]$ is false, or
 - $\text{label}[A], A > \text{label}[B], B$

```
class Bakery implements Lock {  
  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0],  
                        ..., label[n-1]) + 1;  
        while ( $\exists k$  flag[k]  
                && (label[i], i) > (label[k], k));  
    }  
}
```

Mutual Exclusion

- Labels are strictly increasing so
- B must have seen $\text{flag}[A] == \text{false}$

Mutual Exclusion

- Labels are strictly increasing so
- B must have seen $\text{flag}[A] == \text{false}$
- $\text{Labeling}_B \rightarrow \text{read}_B(\text{flag}[A]) \rightarrow$
 $\text{write}_A(\text{flag}[A]) \rightarrow \text{Labeling}_A$

Mutual Exclusion

- Labels are strictly increasing so
- B must have seen $\text{flag}[A] == \text{false}$
- $\text{Labeling}_B \rightarrow \text{read}_B(\text{flag}[A]) \rightarrow$
 $\text{write}_A(\text{flag}[A]) \rightarrow \text{Labeling}_A$
- Which contradicts the assumption that A has an earlier label

Bakery Y2³²K Bug

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i],i) > (label[k],k));  
    }
```


Mutex breaks if label[i] overflows

Timestamps

- Label variable is really a **timestamp**
- Need ability to
 - Read others' timestamps
 - Compare them
 - Generate a **later** timestamp
- Can we do this without overflow?

The Good News

- One can construct a
 - Wait-free (no mutual exclusion)
 - Concurrent
 - Timestamping system
 - That never overflows

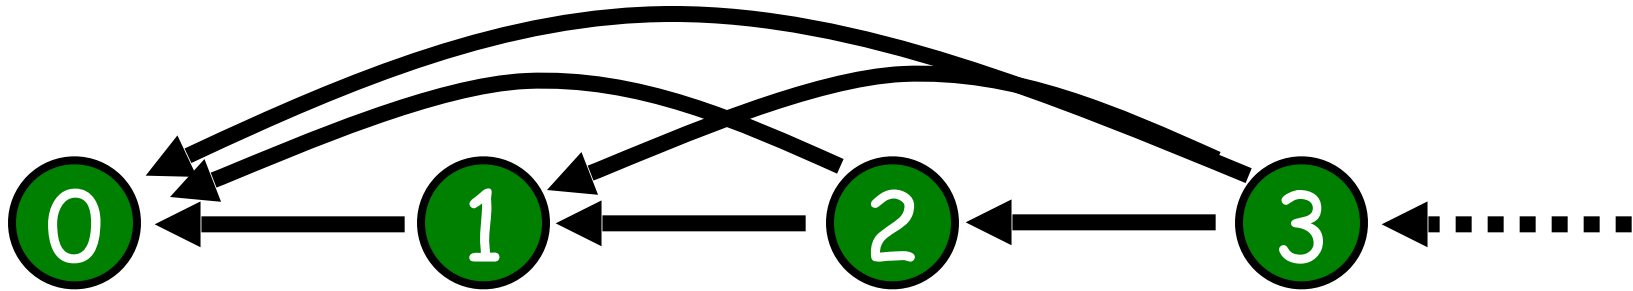
The ~~Good~~ Bad News

- One can construct a
 - Wait-free (no mutual exclusion)
 - Concurrent
 - Timestamping system
 - That never overflows
- This part is hard

Instead ...

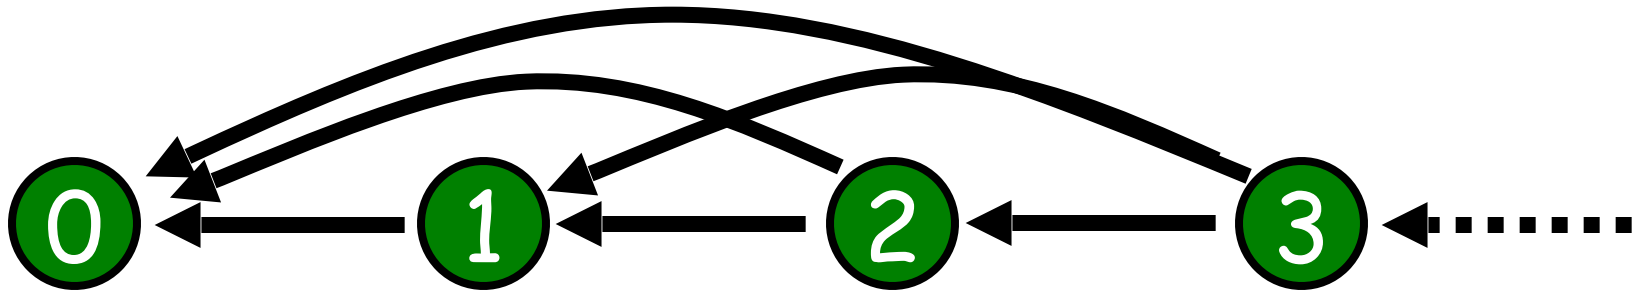
- We construct a Sequential timestamping system
 - Same basic idea
 - But simpler
- Uses mutex to read & write atomically
- No good for building locks
 - But useful anyway

Precedence Graphs



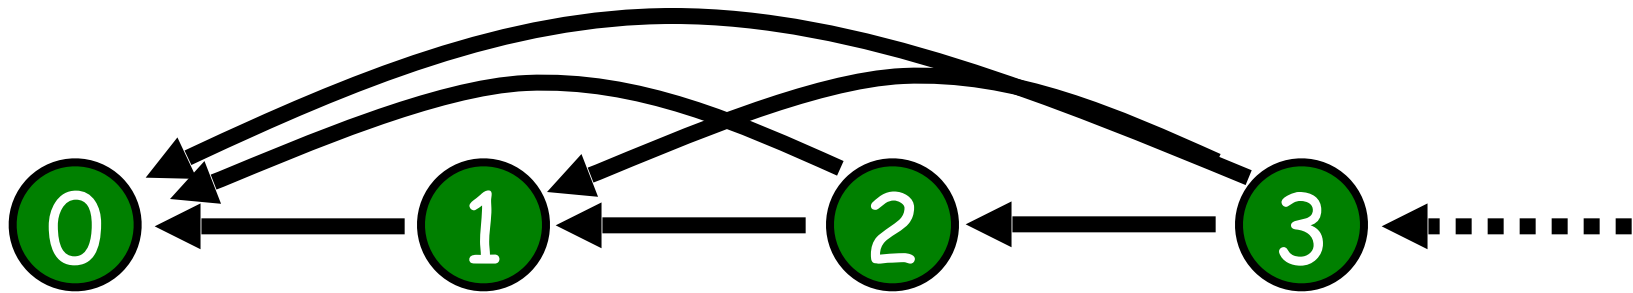
- Timestamps form directed graph
- Edge x to y
 - Means x is later timestamp
 - We say x **dominates** y

Unbounded Counter Precedence Graph

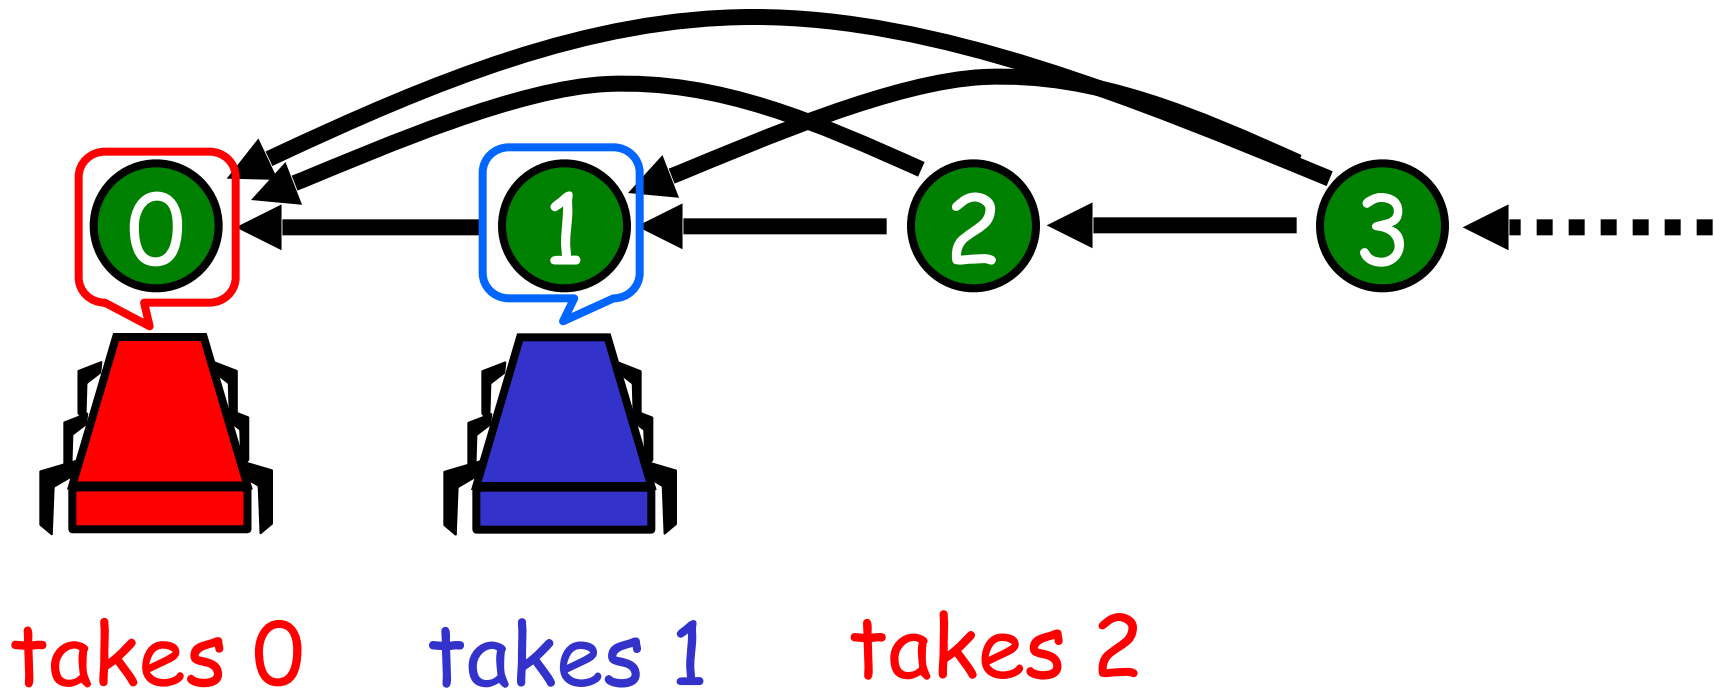


- Timestamping = move tokens on graph
- Atomically
 - read others' tokens
 - move mine
- Ignore tie-breaking for now

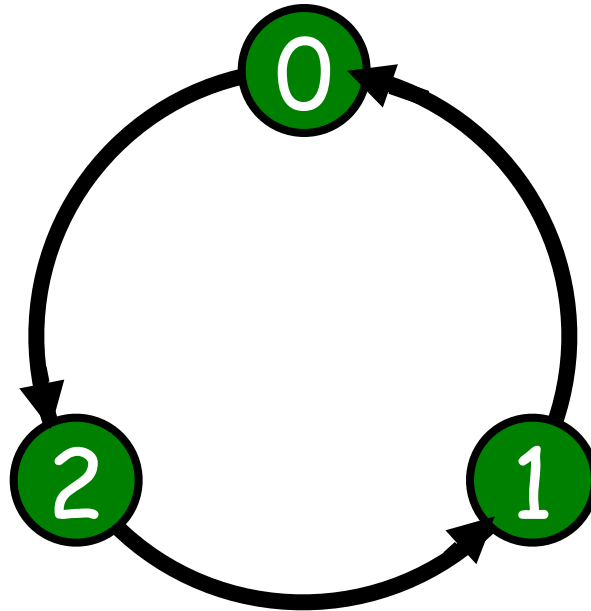
Unbounded Counter Precedence Graph



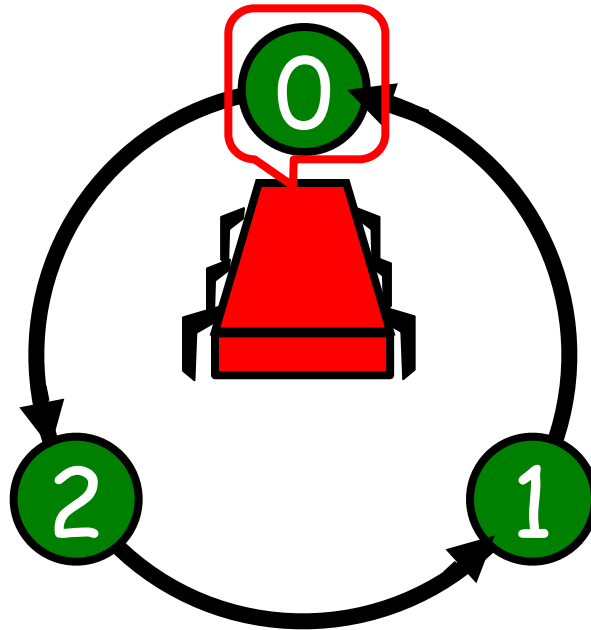
Unbounded Counter Precedence Graph



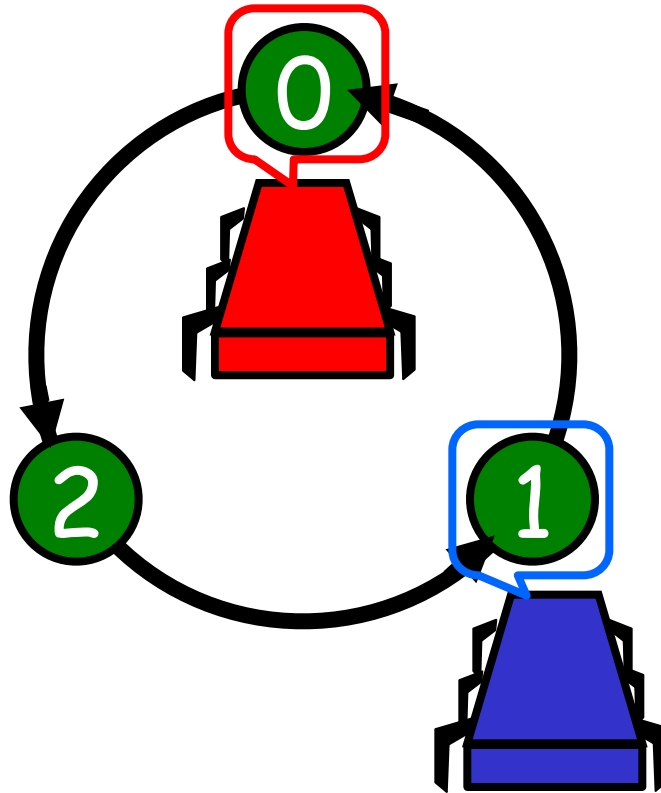
Two-Thread Bounded Precedence Graph



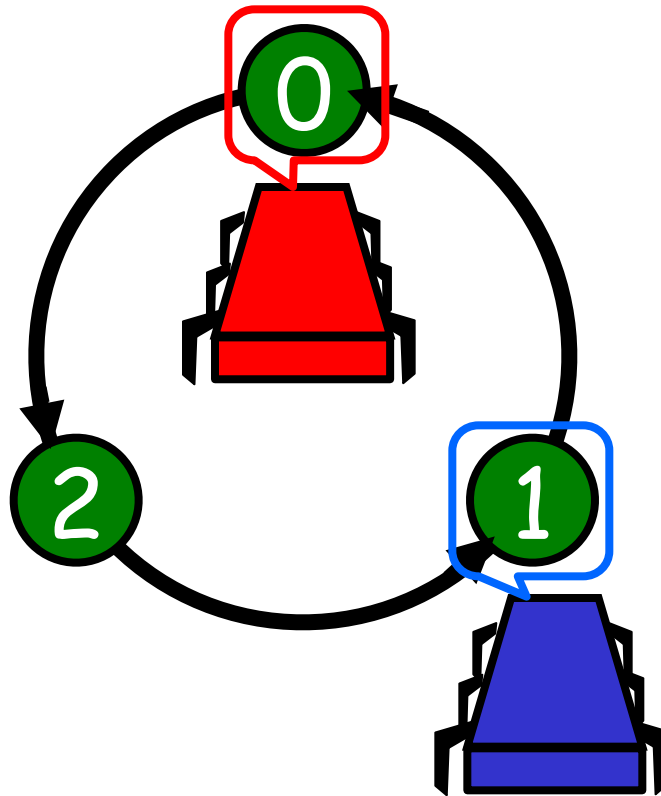
Two-Thread Bounded Precedence Graph



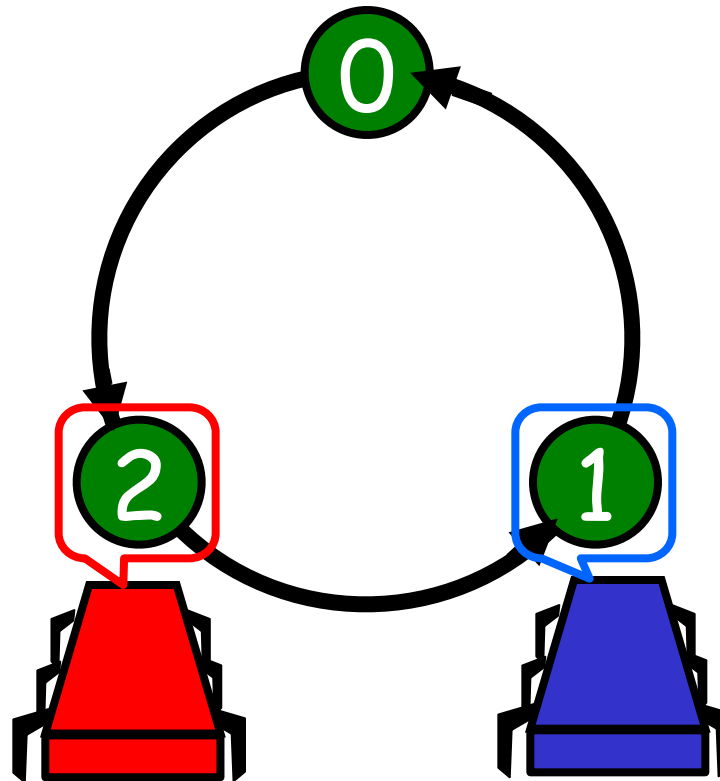
Two-Thread Bounded Precedence Graph



Two-Thread Bounded Precedence Graph

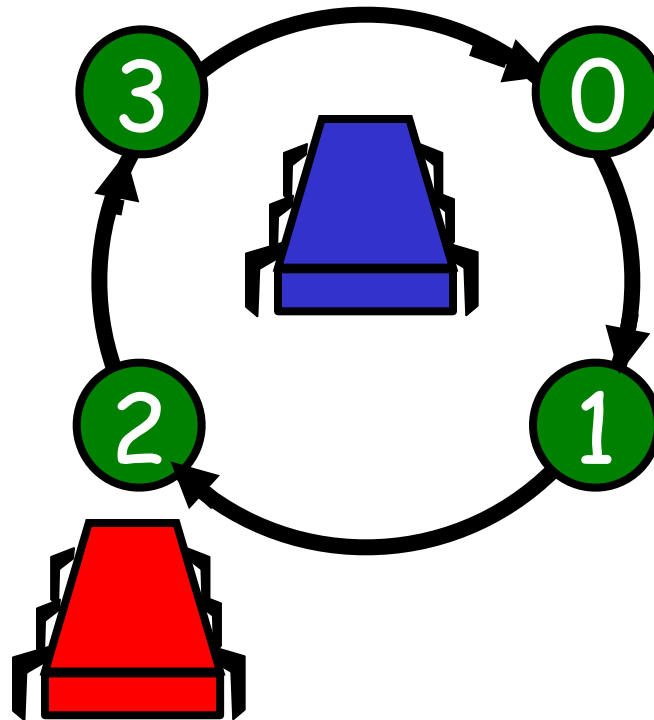


Two-Thread Bounded Precedence Graph T^2

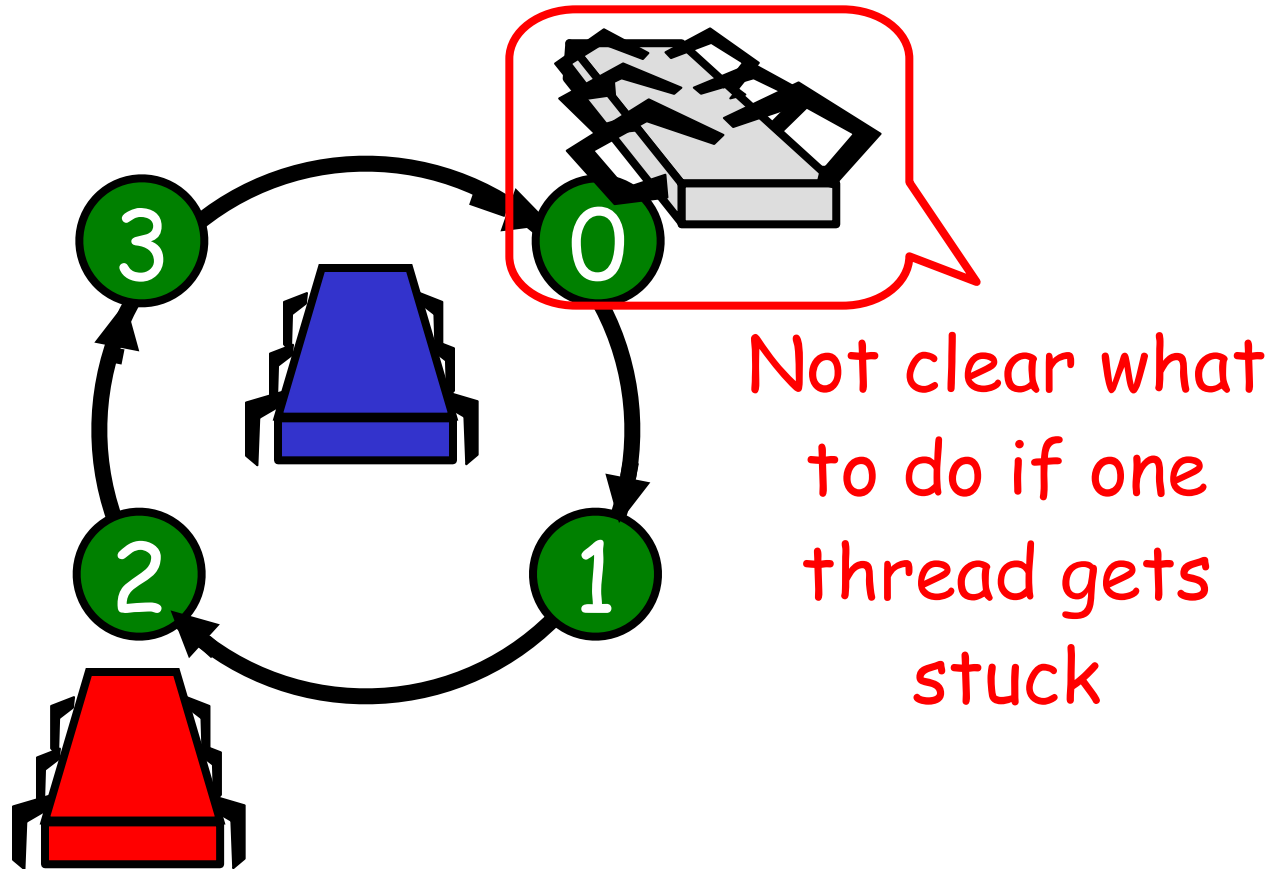


and so on ...

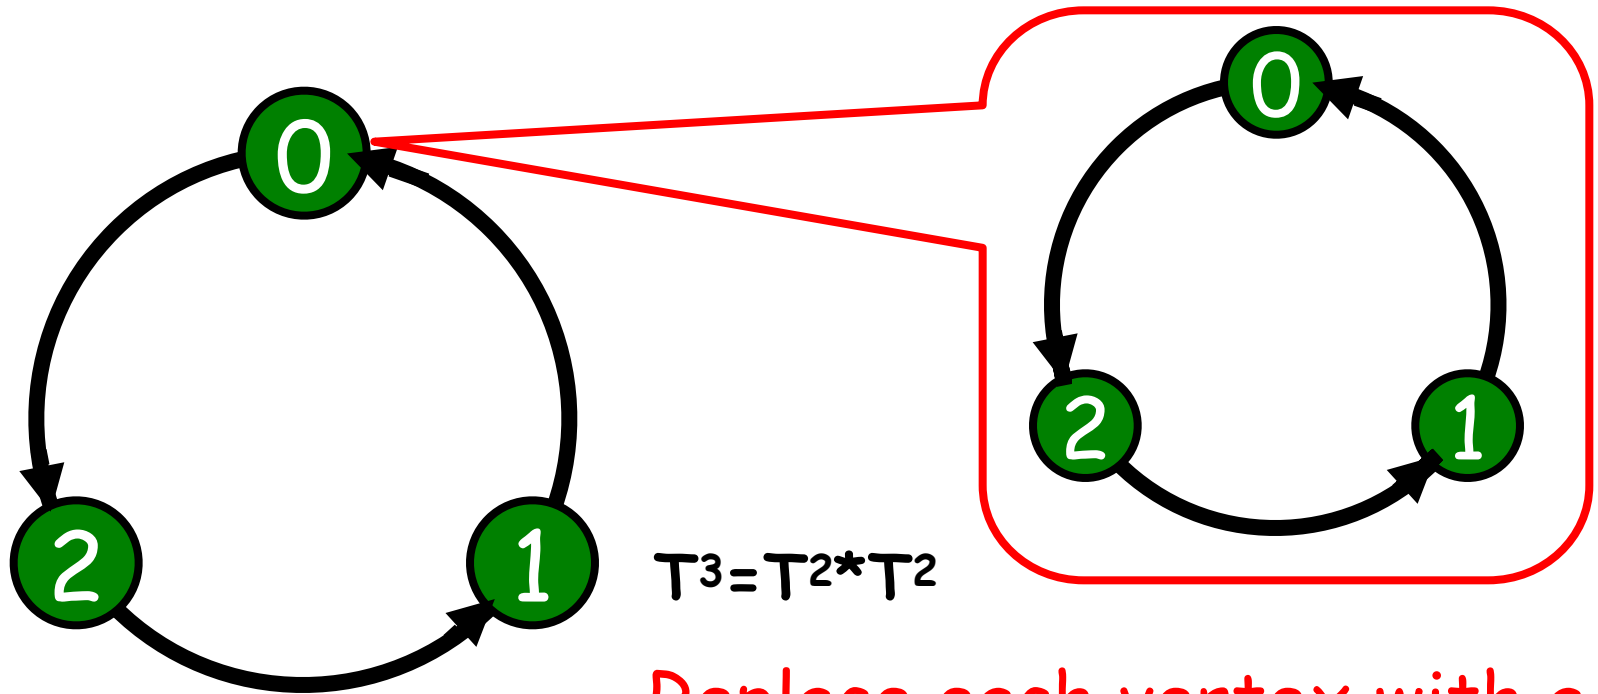
Three-Thread Bounded Precedence Graph?



Three-Thread Bounded Precedence Graph?



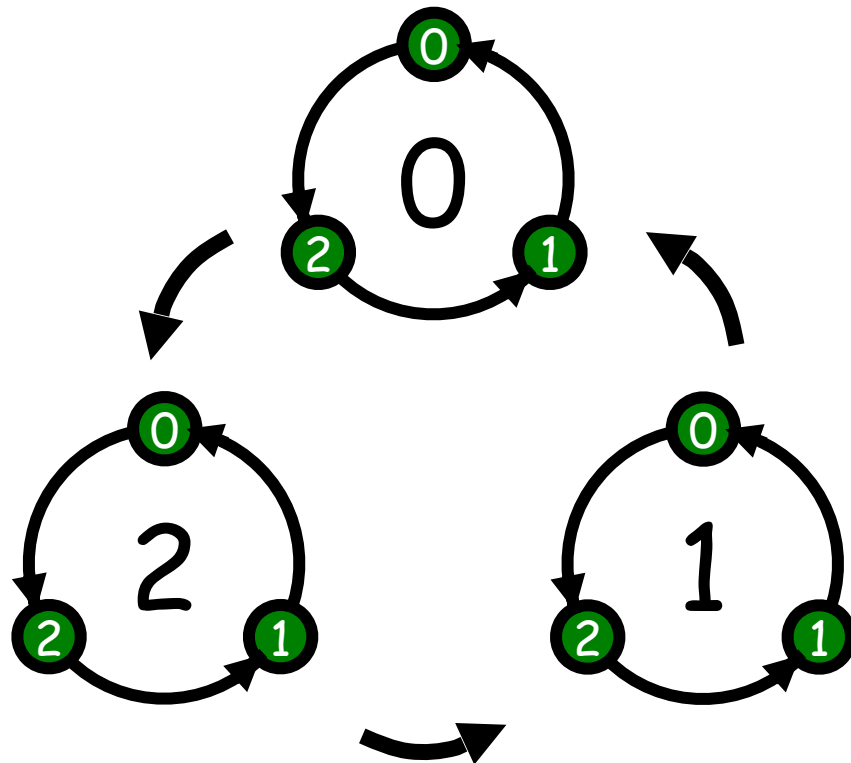
Graph Composition



Replace each vertex with a
copy of the graph

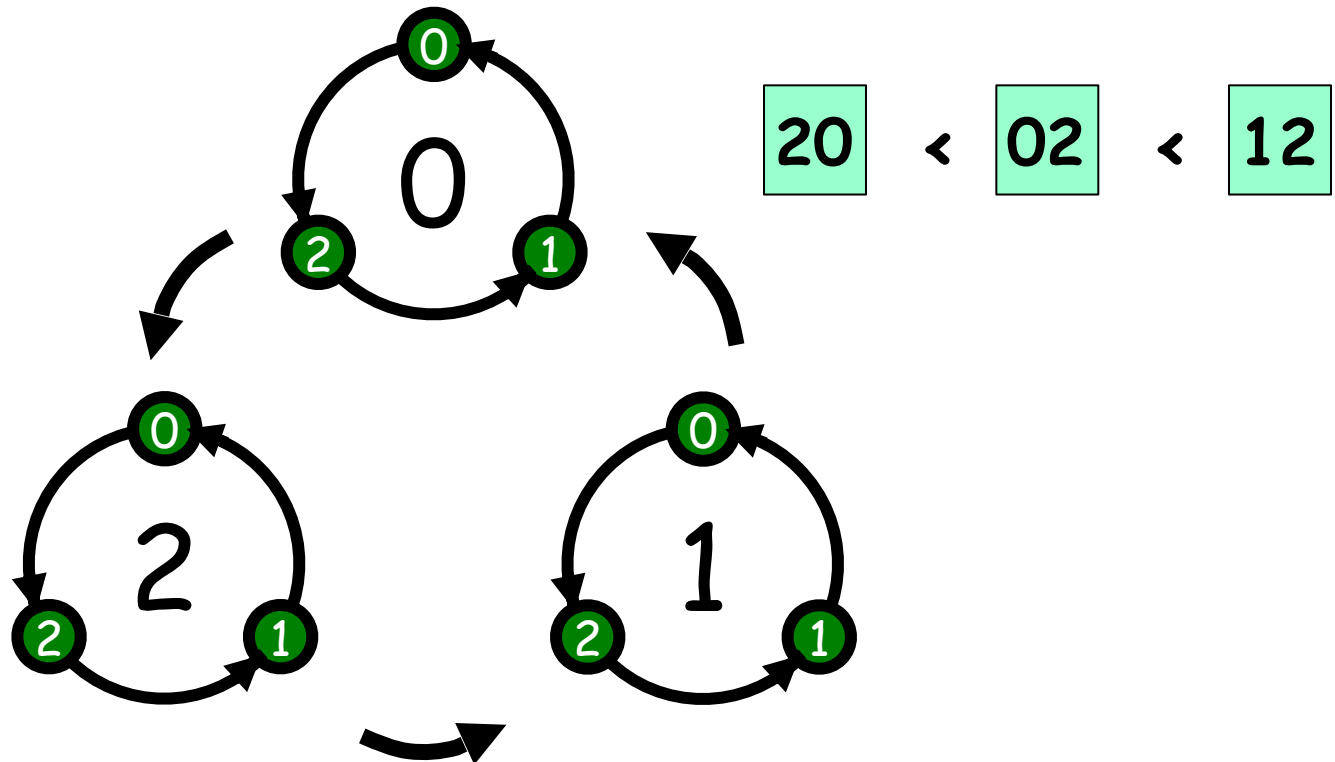
Three-Thread Bounded Precedence Graph

T_3



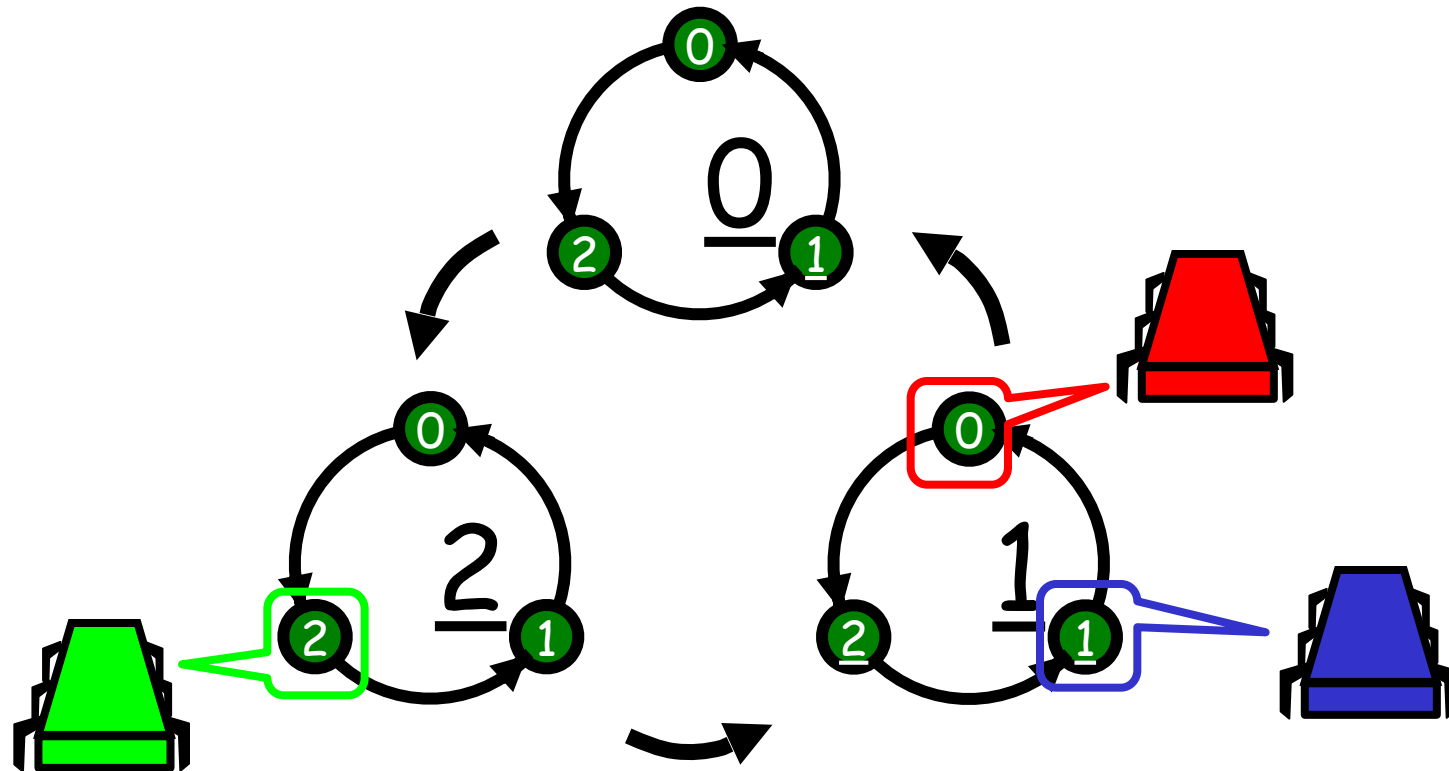
Three-Thread Bounded Precedence Graph

T_3



Three-Thread Bounded Precedence Graph

T_3



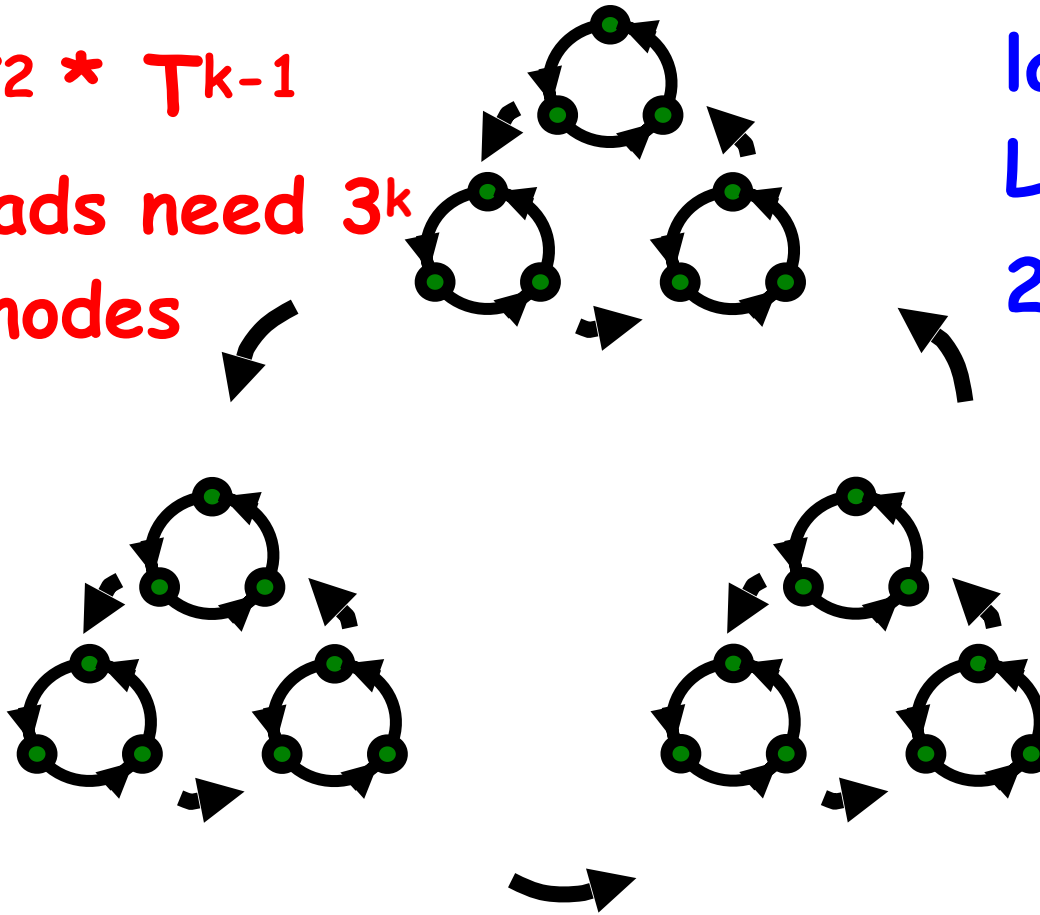
and so on...

In General

$$T_k = T_2 * T_{k-1}$$

K threads need 3^k
nodes

$$\text{label size} = \log_2(3^k) = 2n$$



Deep Philosophical Question

- The Bakery Algorithm is
 - Succinct,
 - Elegant, and
 - Fair.
- Q: So why isn't it practical?
- A: Well, you have to read **N** distinct variables

Shared Memory

- Shared read/write memory locations called **Registers** (historical reasons)
- Come in different flavors
 - Multi-Reader-Single-Writer (**Flag**[])
 - Multi-Reader-Multi-Writer (**Victim**[])
 - Not that interesting: SRMW and SRSW

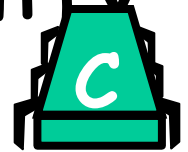
Theorem

At least N MRSW (multi-reader/
single-writer) registers are needed
to solve deadlock-free mutual
exclusion.

N registers like `Flag[]...`

Proving Algorithmic Impossibility

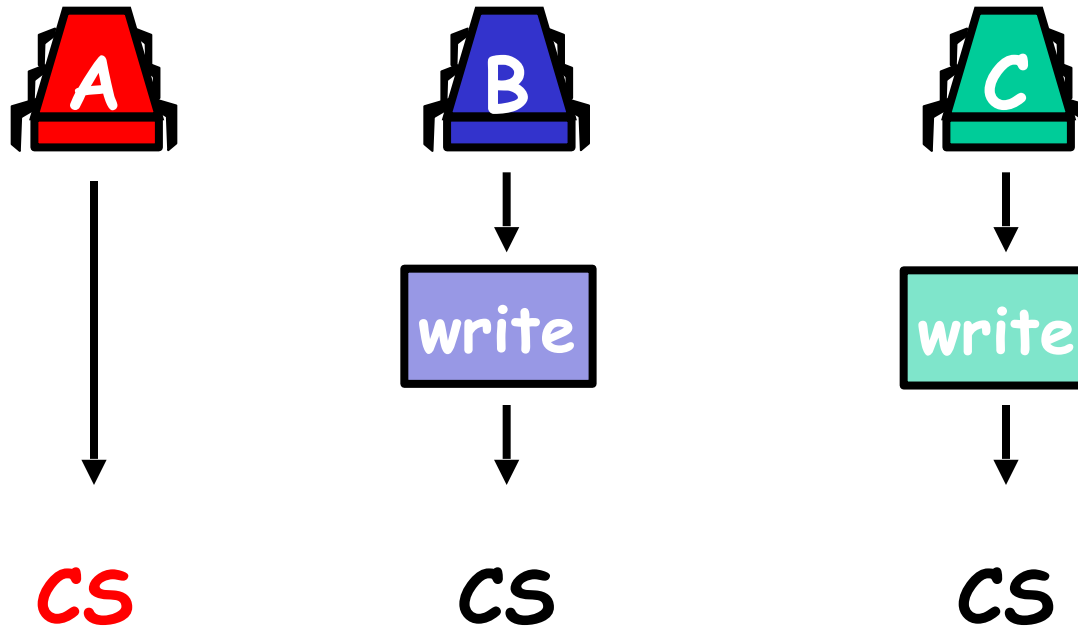
- To show no algorithm exists:
 - assume by way of contradiction one does,
 - show a **bad execution** that violates properties:
 - in our case assume an alg for deadlock free mutual exclusion using $< N$ registers



CS

Proof: Need N-MRSW Registers

Each thread must write to some register



...can't tell whether **A** is in critical
section

Upper Bound

- Bakery algorithm
 - Uses $2N$ MRSW registers
- So the bound is (pretty) tight
- But what if we use MRMW registers?
 - Like `victim[]` ?

Bad News Theorem

At least N MRMW multi-reader/
multi-writer registers are needed
to solve deadlock-free mutual
exclusion.

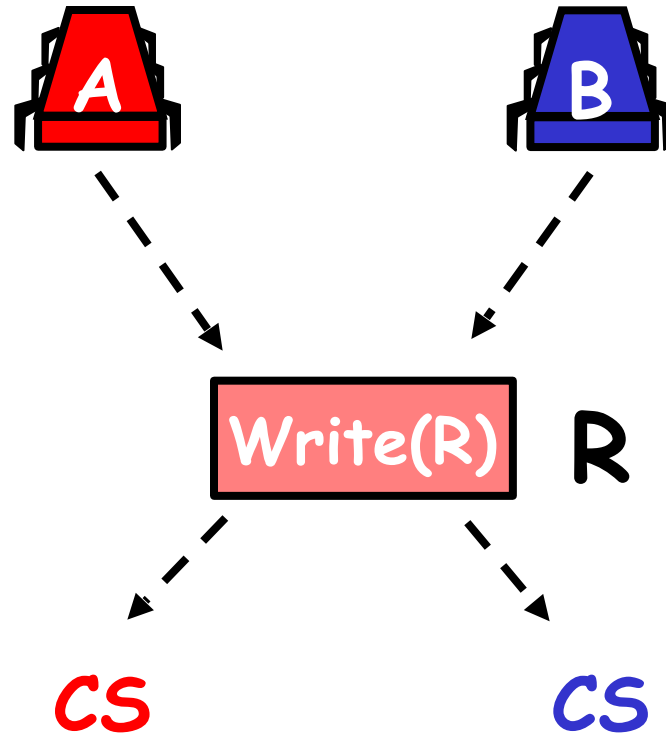
(So multiple writers don't help)

Theorem (First 2-Threads)

Theorem: Deadlock-free mutual exclusion for 2 threads requires at least 2 multi-reader multi-writer registers

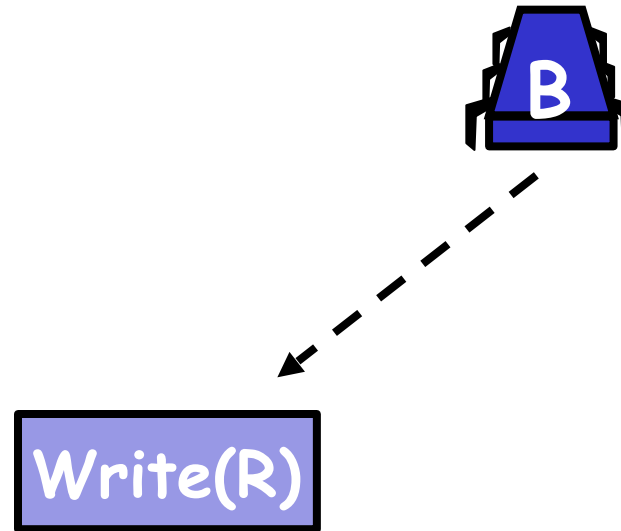
Proof: assume one register suffices and derive a contradiction

Two Thread Execution



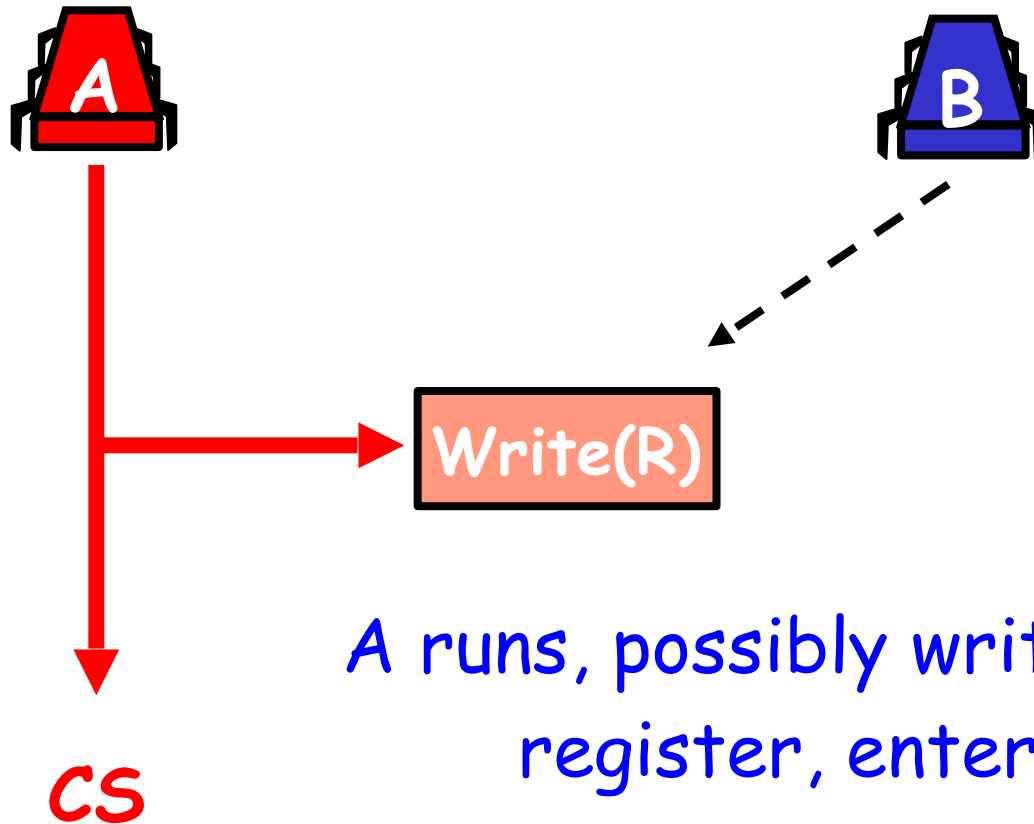
- Threads run, reading and writing R
- Deadlock free so at least one gets in

Covering State for One Register Always Exists

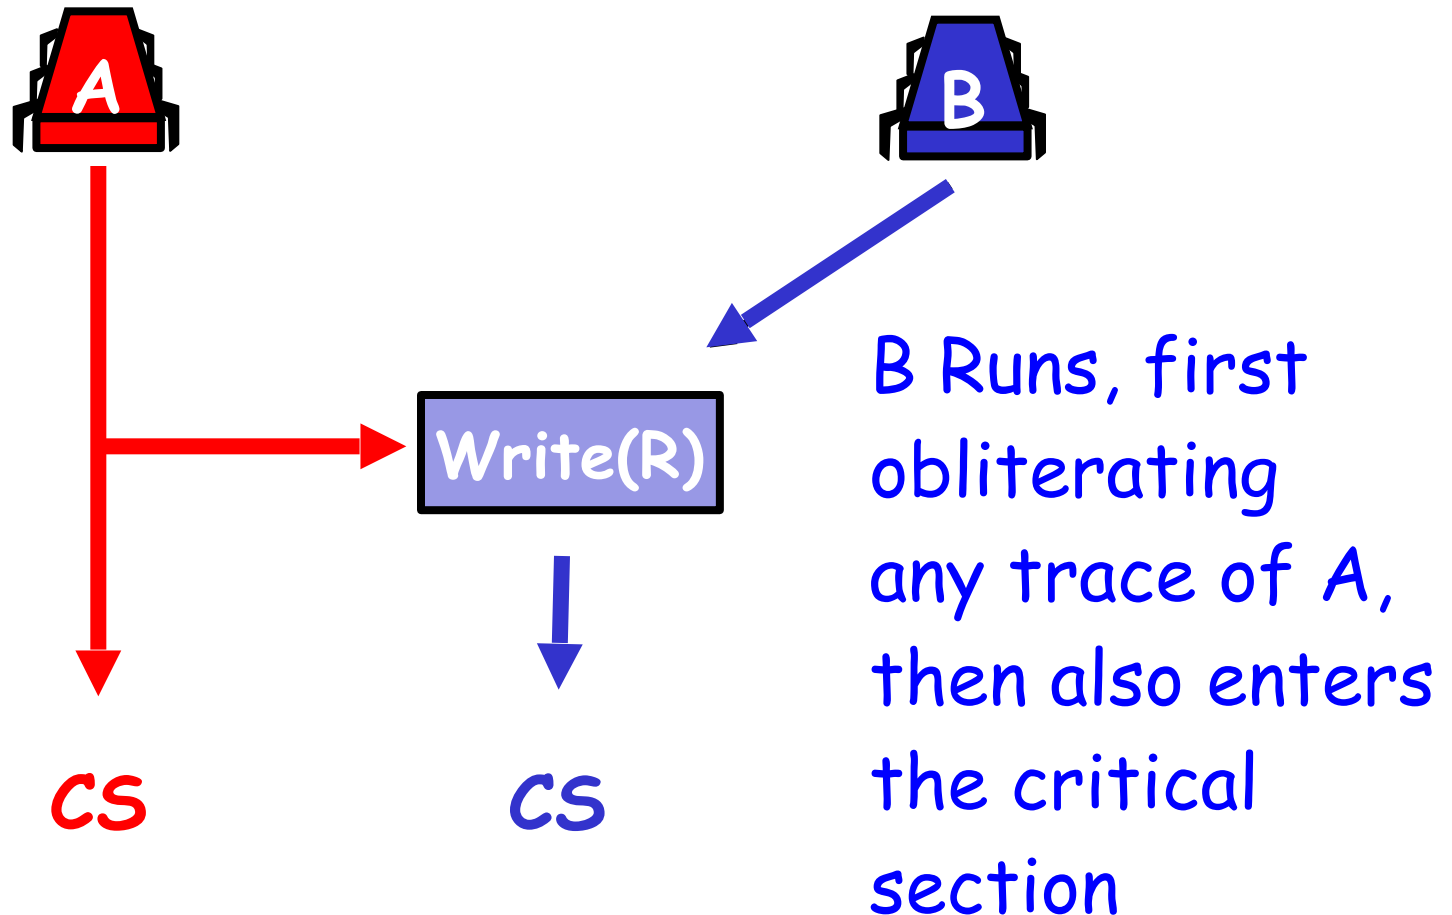


In any protocol B has to write to the register before entering CS, so stop it just before

Proof: Assume Cover of 1



Proof: Assume Cover of 1



Theorem

Deadlock-free mutual exclusion for 3 threads requires at least 3 multi-reader multi-writer registers